

# Example 1:

## Creating a Cassandra Cluster

# What will we learn in this example?


**Bird's eye view of a cluster**

**Create a cassandra cluster on your local machine**

Bird's eye view of a cluster

**Cassandra** is a distributed,  
decentralized database

one ring is called  
data centre




This cluster has 1  
datacenter


**Cassandra Cluster**

*Bird's eye view of a cluster*

## **Multiple Datacenters**



**Datacenter1**





**Datacenter2**



Bird's eye view of a cluster

**datacenters can be at different  
geographical locations**







Bird's eye view of a cluster

**But why would we do that?**

**Users connect to datacenters based  
on geographic location**

**This way we can ensure faster  
performance for end users.**



# What will we learn in this example?

~~Bird's eye view of a cluster~~

Create a cassandra cluster on your local machine

Create a cassandra cluster on your local machine

We will use the **Cassandra Cluster Manager (ccm)**

**note that ccm cannot be used on production**

**Its made for 2 purposes**

- 1. To learn cassandra**
- 2. To test an application using cassandra on your local environment**



# Create a cassandra cluster on your local machine

**ccm** has a list of commands to manage the cluster

You can run **ccm help** to see the full list of commands

# Create a cassandra cluster on your local machine

## On its start, cassandra opens 3 ports on a node

9160

thrift based clients connect to  
cassandra using this port number

9042

binary protocol based clients connect  
to cassandra using this port number  
eg CQL

7000

Nodes communicate with each other  
using this port number

Create a cassandra cluster on your local machine

**We create a cluster on 1 machine  
(local)**

**Create a network interface on loopback  
addresses (127.0.0.\*)**

**In windows and unix, the network  
interface aliasing is done internally by ccm**

**For macOS, we need to do it manually**



# Create a cassandra cluster on your local machine

If you have **n** nodes in your cluster  
then run the following command for all **n** of them

```
sudo ifconfig lo0 alias 127.0.0.2
```

```
sudo ifconfig lo0 alias 127.0.0.3
```

```
.
```

```
.
```

```
sudo ifconfig lo0 alias 127.0.0.n
```

**Run these commands BEFORE you run  
the command to create your cluster**

Create a cassandra cluster on your local machine

Let's create a cluster for "easybuy"

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

Command to create  
the cluster

Create a cassandra cluster on your local machine

Let's create a cluster for "easybuy"

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

Name of the cluster  
to be created



Create a cassandra cluster on your local machine

Let's create a cluster for "easybuy"

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

Options for create  
command

You can see the full list of options for  
create command by running

```
ccm create help
```

# Create a cassandra cluster on your local machine

## Lets go through the options

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

version of  
Cassandra installed

To know the version of Cassandra run the following command on the terminal

```
cassandra -v
```

# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

this option  
populates the cluster  
with the number of  
nodes passed



# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

this option populates the cluster with the  
number of nodes passed

**We can create the cluster even without this option**

**In that case we will have an empty cluster without  
any nodes**

# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

this option populates the cluster with the  
number of nodes passed

We can create the cluster even without this option

**Run the following command to populate cluster  
with 5 nodes**

```
ccm populate -n 5
```

# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

**this option starts the cassandra cluster**  
applications can now connect to the cluster



# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

Again, the cluster can be created without this option

If you are not using this option then you start the cluster with the following command

```
ccm start
```

run the command after  
you have populated the cluster

# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

**What is an  
authenticator?**

**we set the authenticator to  
PasswordAuthenticator**

**Authenticator is the class that  
authenticates clients/users  
connecting to cassandra**

# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

default authenticator is  
**AllowAllAuthenticator**

No user/password required

**Anyone** can connect to the cluster

**This is not very secure - pretty obvious**

# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

We want that only the application  
and our team should be able to  
connect to the cluster

We can ensure this by using the  
PasswordAuthenticator



# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

To connect to cassandra,  
**PasswordAuthenticator** forces the user to  
provide a valid username/password

# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

**default** username/password for connecting to the cluster is "cassandra"

**We can configure our own username/  
password for our cluster**

# Create a cassandra cluster on your local machine

```
ccm create easybuy -v 3.7 -n 5 -s --pwd-auth
```

If the cluster is created successfully then the output is :

```
Current cluster is now: easybuy
```

# What will we learn in this example?

~~Bird eye view of a cluster~~

~~Create a cassandra cluster on your local machine~~



# Example 2: Exploring CCM

Let's see the basic  
commands of CCM

# Basic CCM commands

```
ccm list
```

This lists all the existing clusters

output

```
*easybuy
```

\* denotes that this is the  
current cluster

# Basic CCM commands

```
ccm status
```

This displays status of the nodes  
in the current cluster

output

```
Cluster: 'easybuy'
```

```
-----
```

```
node1: UP
```

```
node3: UP
```

```
node2: UP
```

```
node5: UP
```

```
node4: UP
```

**UP** - node is ready to take  
requests

**DOWN** - node is not ready  
to take requests



# Basic CCM commands

```
ccm stop
```

**Stops all the nodes in the cluster**

**After we have stopped the cluster, the output of the `ccm status` command will be:**

```
Cluster: 'easybuy'
```

```
-----
```

```
node1: DOWN
```

```
node3: DOWN
```

```
node2: DOWN
```

```
node5: DOWN
```

```
node4: DOWN
```

# Basic CCM commands

```
ccm node1 show
```

Displays information of node named  
"node1"

output

```
node1: UP
      cluster=easybuy
      auto_bootstrap=False
      thrift=('127.0.0.1', 9160)
      binary=('127.0.0.1', 9042)
      storage=('127.0.0.1', 7000)
      jmx_port=7100
      remote_debug_port=0
      byteman_port=0
      initial_token=-9223372036854775808
      pid=3028
```

# Basic CCM commands

Lets understand the fields of the output

```
ccm node1 show
```

```
node1: UP
```

```
cluster=easybuy
auto_bootstrap=False
thrift=('127.0.0.1', 9160)
binary=('127.0.0.1', 9042)
storage=('127.0.0.1', 7000)
jmx_port=7100
remote_debug_port=0
byteman_port=0
initial_token=-9223372036854775808
pid=3028
```

Name of the node and the status



# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
```

```
cluster=easybuy
```

```
auto_bootstrap=False
```

```
thrift=('127.0.0.1', 9160)
```

```
binary=('127.0.0.1', 9042)
```

```
storage=('127.0.0.1', 7000)
```

```
jmx_port=7100
```

```
remote_debug_port=0
```

```
byteman_port=0
```

```
initial_token=-9223372036854775808
```

```
pid=3028
```

**node1 belongs to this cluster**



# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
cluster=easybuy
auto_bootstrap=False
thrift=('127.0.0.1', 9160)
binary=('127.0.0.1', 9042)
storage=('127.0.0.1', 7000)
jmx_port=7100
remote_debug_port=0
byteman_port=0
initial_token=-9223372036854775808
pid=3028
```

new node will  
automatically  
migrate data from  
the cluster if the flag  
is true

default value for the  
flag is false

# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
cluster=easybuy
auto_bootstrap=False
thrift=('127.0.0.1', 9160)
binary=('127.0.0.1', 9042)
storage=('127.0.0.1', 7000)
jmx_port=7100
remote_debug_port=0
byteman_port=0
initial_token=-9223372036854775808
pid=3028
```

For a new cluster  
without any data, set  
the value False

If you are adding a node  
to a cluster with data

set the flag to true

**ONLY** after the new node is  
properly configured

# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
      cluster=easybuy
      auto_bootstrap=False
      thrift=('127.0.0.1', 9160)
      binary=('127.0.0.1', 9042)
      storage=('127.0.0.1', 7000)
      jmx_port=7100
      remote_debug_port=0
      byteman_port=0
      initial_token=-9223372036854775808
      pid=3028
```

**As we had seen before  
this port is used for  
connecting via thrift**



# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
cluster=easybuy
auto_bootstrap=False
thrift=('127.0.0.1', 9160)
binary=('127.0.0.1', 9042)
storage=('127.0.0.1', 7000)
jmx_port=7100
remote_debug_port=0
byteman_port=0
initial_token=-9223372036854775808
pid=3028
```

this port is used for  
connecting from cqlsh  
or any client using  
binary protocol



# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
cluster=easybuy
auto_bootstrap=False
thrift=('127.0.0.1', 9160)
binary=('127.0.0.1', 9042)
storage=('127.0.0.1', 7000)
jmx_port=7100
remote_debug_port=0
byteman_port=0
initial_token=-9223372036854775808
pid=3028
```

this port is used for  
inter-node  
communication

# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
cluster=easybuy
auto_bootstrap=False
thrift=('127.0.0.1', 9160)
binary=('127.0.0.1', 9042)
storage=('127.0.0.1', 7000)
jmx_port=7100
remote_debug_port=0
byteman_port=0
initial_token=-9223372036854775808
pid=3028
```

Port to monitor the  
cluster using java  
based monitoring  
tools **e.g. jconsole**

# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
  cluster=easybuy
  auto_bootstrap=False
  thrift=('127.0.0.1', 9160)
  binary=('127.0.0.1', 9042)
  storage=('127.0.0.1', 7000)
  imx_port=7100
  remote_debug_port=0
  byteman_port=0
  initial_token=-9223372036854775808
  pid=3028
```

**Port for remote  
debugging from the  
application. By default  
it is not set**



# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
  cluster=easybuy
  auto_bootstrap=False
  thrift=('127.0.0.1', 9160)
  binary=('127.0.0.1', 9042)
  storage=('127.0.0.1', 7000)
  jmx_port=7100
  remote debug port=0
  byteman_port=0
  initial_token=-9223372036854775808
  pid=3028
```

**Byteman tests/  
BMUnits use this port  
for connecting to  
cassandra cluster**

# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
  cluster=easybuy
  auto_bootstrap=False
  thrift=('127.0.0.1', 9160)
  binary=('127.0.0.1', 9042)
  storage=('127.0.0.1', 7000)
  jmx_port=7100
  remote_debug_port=0
  hysteresis_port=0
  initial_token=-9223372036854775808
  pid=3028
```

hash-code of the first  
primary key value  
stored on this node.

we will understand  
tokens in detail later  
on when we talk about  
the Cassandra  
architecture



# Basic CCM commands

```
ccm node1 show
```

```
node1: UP
  cluster=easybuy
  auto_bootstrap=False
  thrift=('127.0.0.1', 9160)
  binary=('127.0.0.1', 9042)
  storage=('127.0.0.1', 7000)
  jmx_port=7100
  remote_debug_port=0
  byteman_port=0
  initial_token=-9223372036854775808
  pid=3028
```

**processId of cassandra  
on the node**

# Basic CCM commands

```
ccm remove
```

This destroys the current cluster  
and all data in the cluster is  
**deleted permanently**

**Use this command cautiously**

# We saw the basic commands of CCM to:

check the status of the cluster

see the configuration on a node

start the cluster

stop the cluster

These should be enough to get us started with using Cassandra