

# Cassandra Architecture is inspired by




# Cassandra

Log Structured ColumnFamily Data -  
Based on Google Big table

Partitioning and Replication -  
Based on Amazon DynamoDB

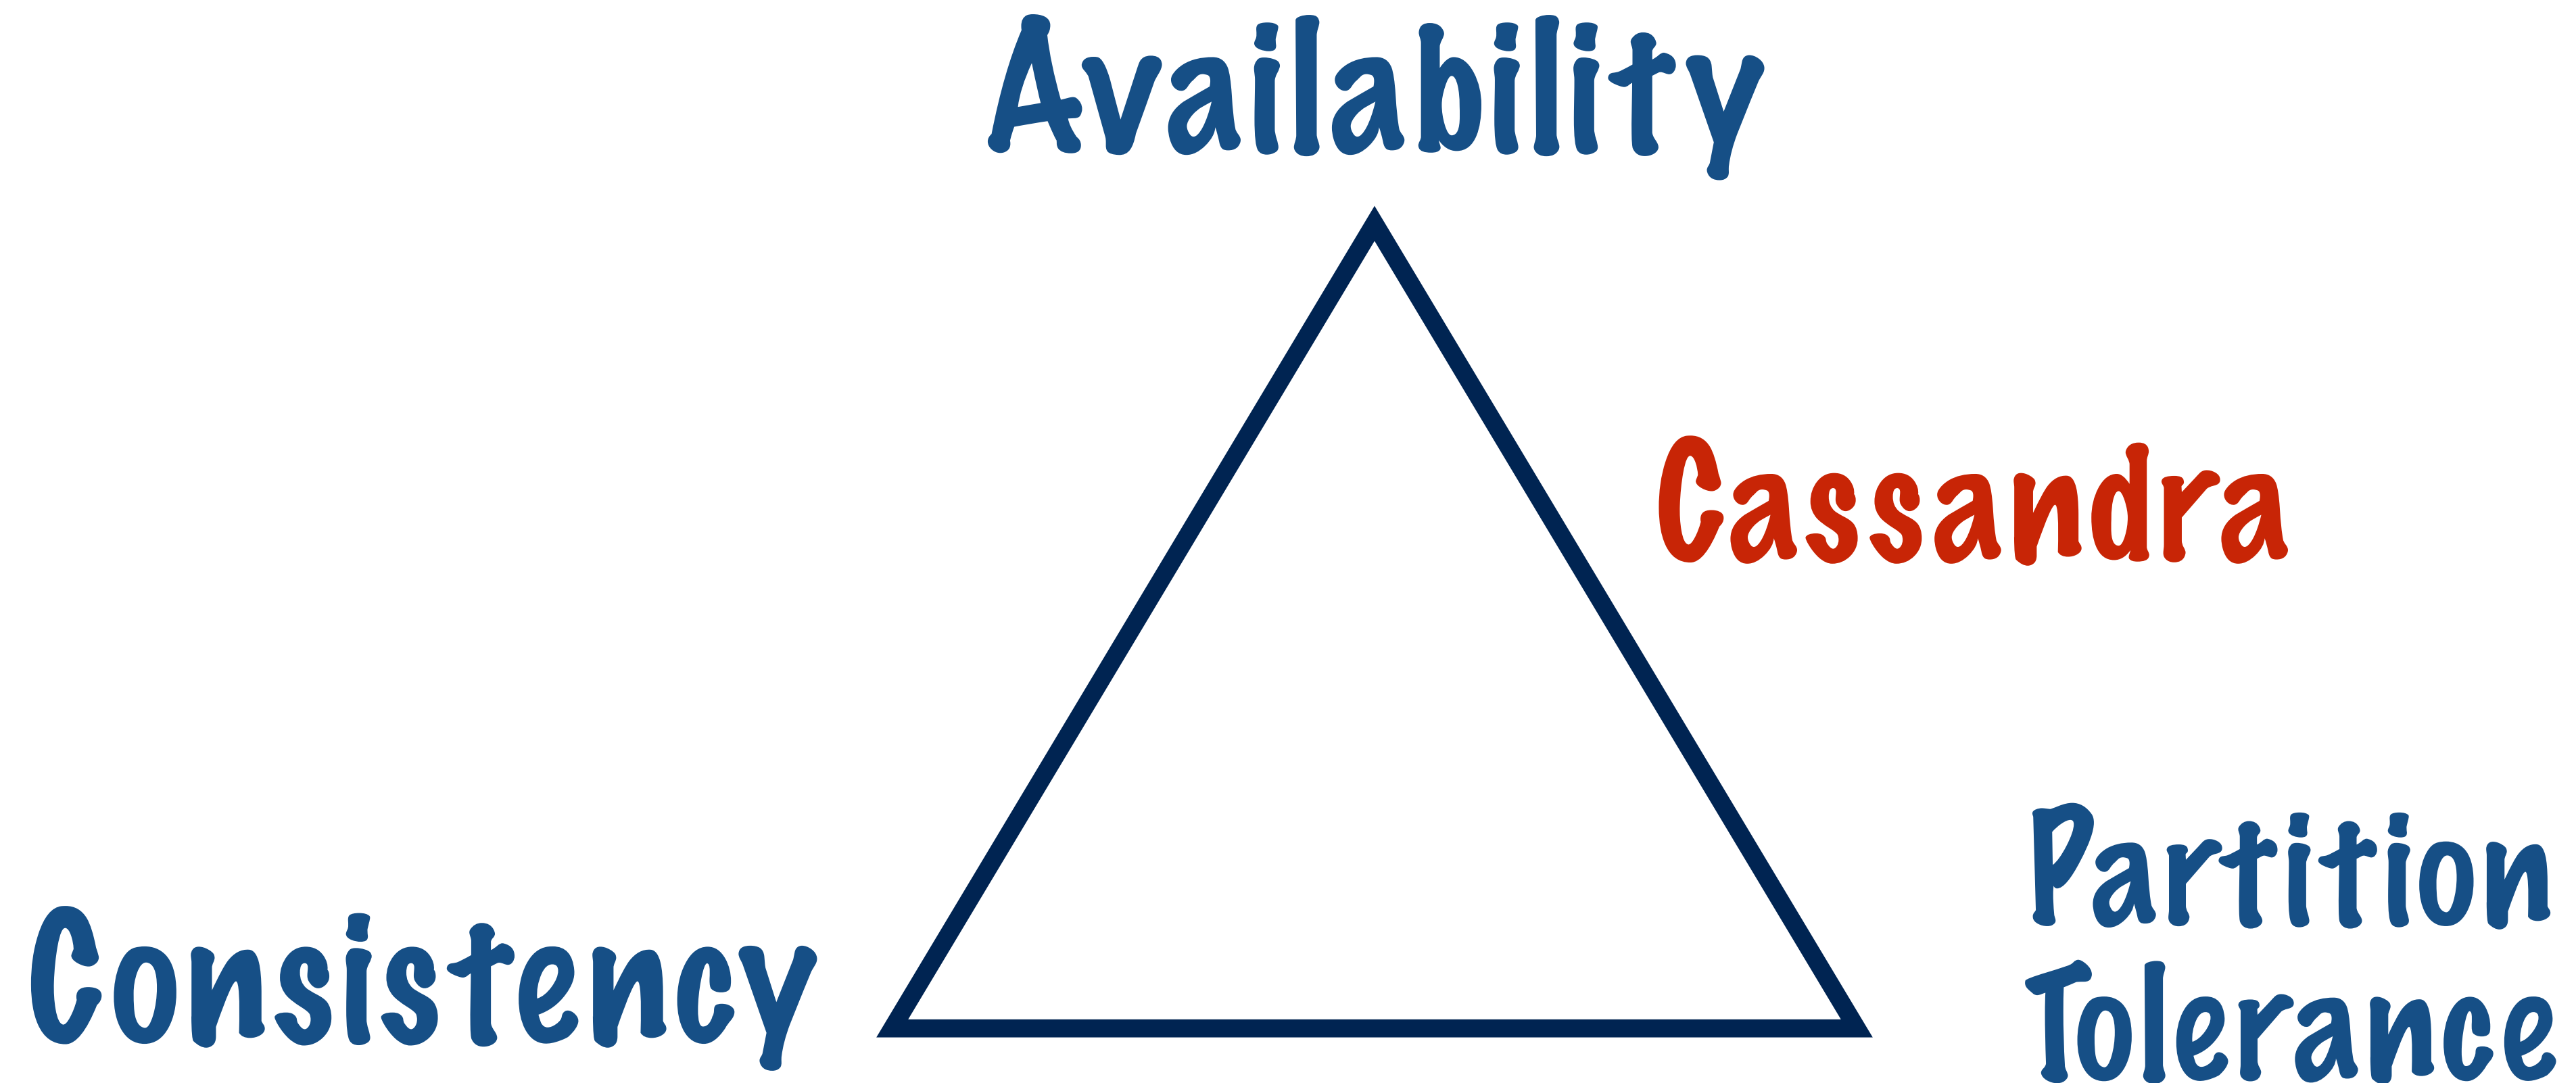
# Cassandra

According to CAP theorem, a distributed system can only fulfil 2 of the following 3 properties



# Cassandra

Cassandra is available and partition tolerant.  
But it trades off **consistency** with **performance**



Let's understand  
consistency in Cassandra

# CONSISTENCY

Consistency ensures that data read from any node in the cluster is the same i.e consistent

# CONSISTENCY

Can configure separate  
consistency levels for  
**READ** and **WRITE**

# CONSISTENCY

## WRITE

number of replica nodes on which  
the **write must succeed**  
before returning success  
to the client

## READ

number of replica nodes  
to **check** before returning  
data to client



# CONSISTENCY

WRITE

ONE

ALL

QUORUM

LOCAL\_QUORUM

# CONSISTENCY

WRITE

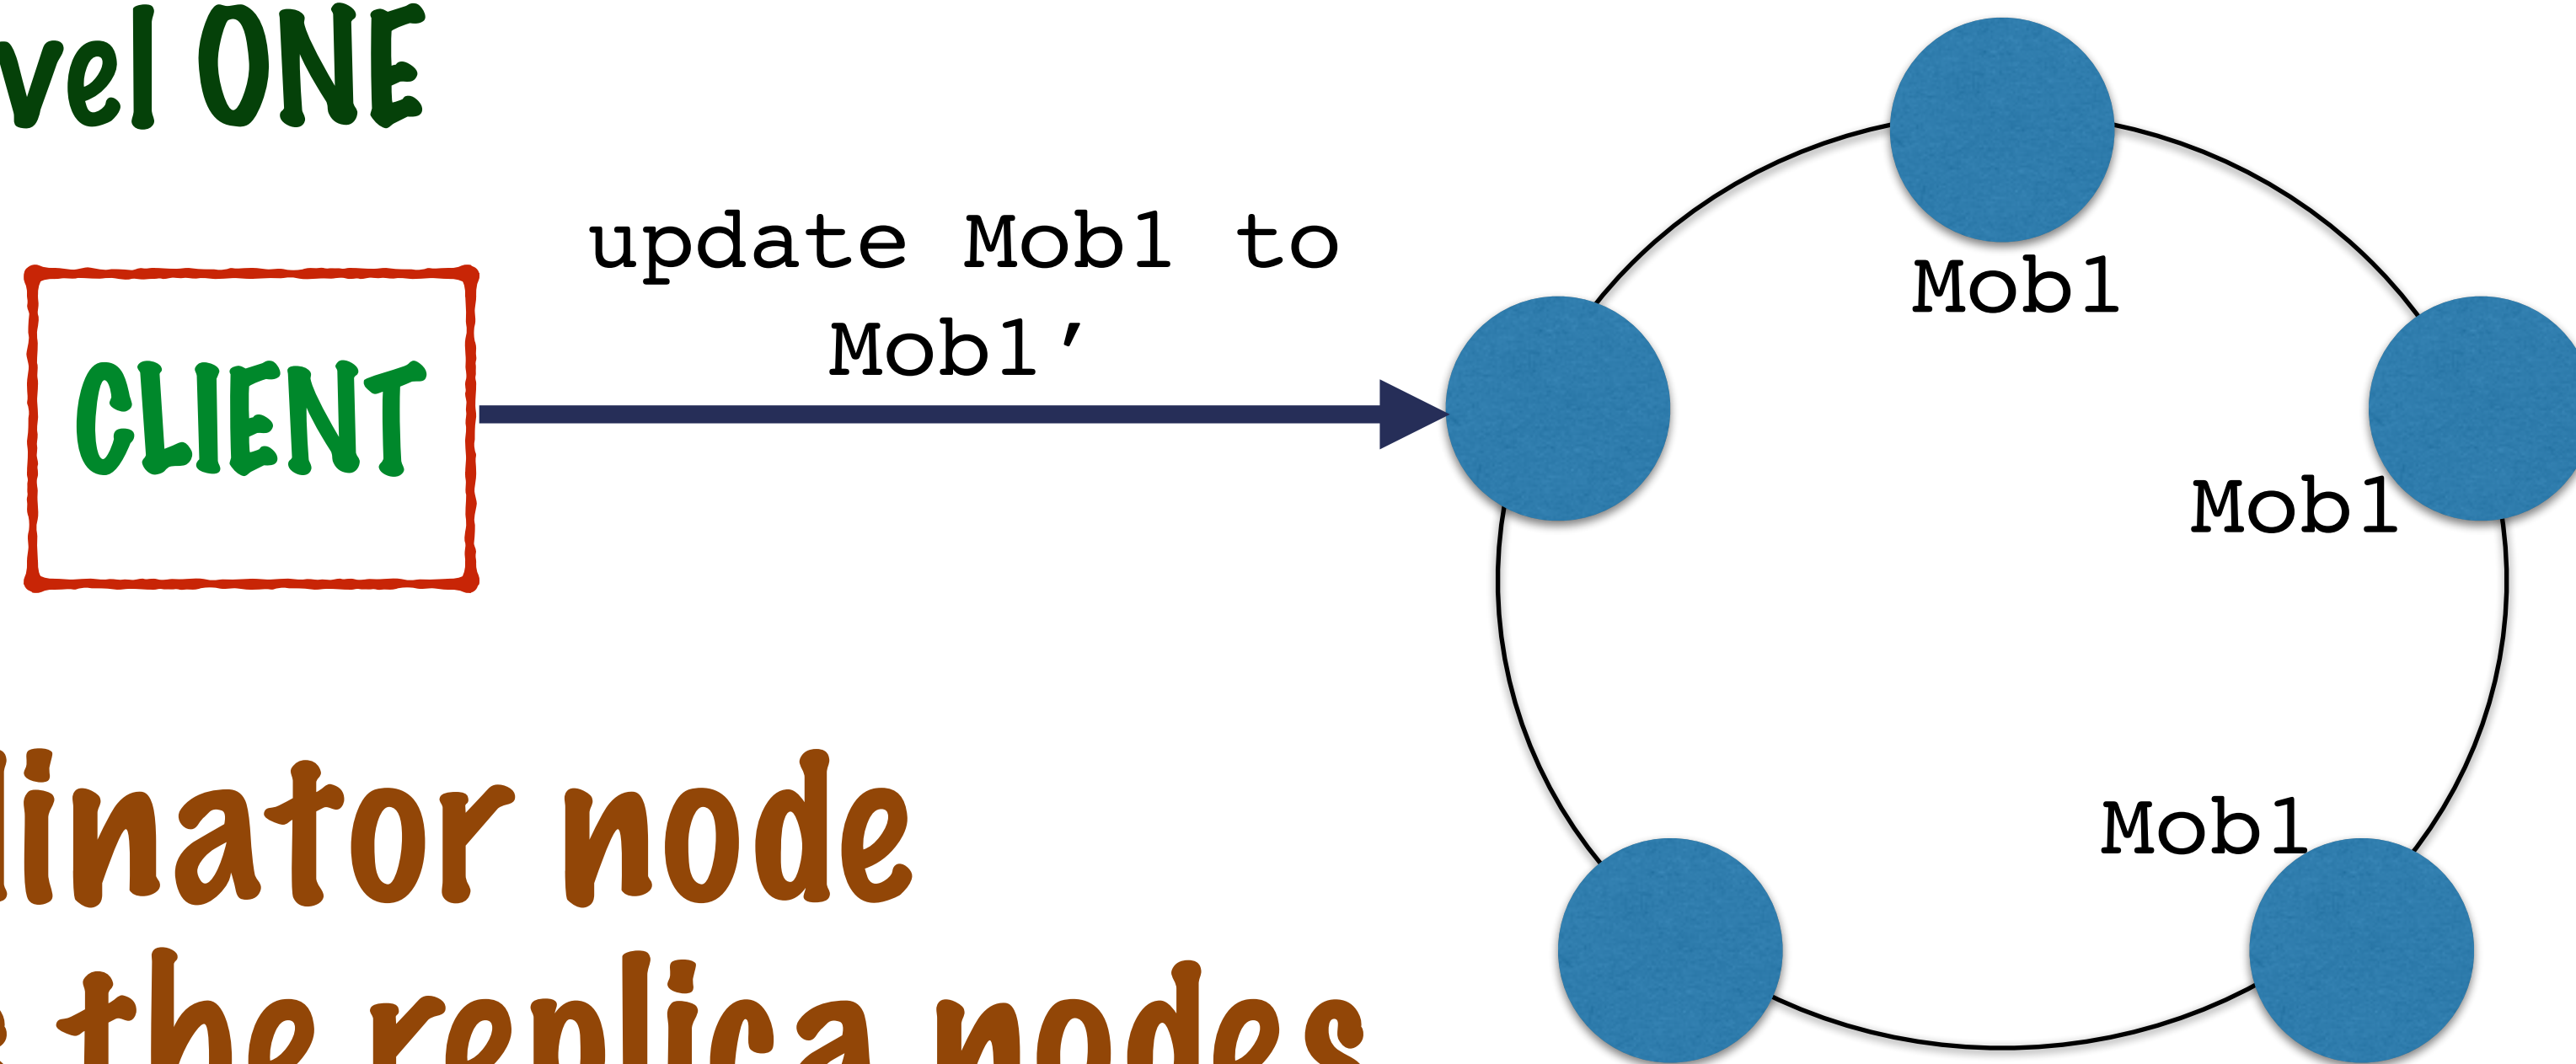
ONE

**Only one** replica needs to be updated and then the write operation returns a success

# CONSISTENCY

## WRITE

Consistency Level ONE




coordinator node  
determines the replica nodes  
with the use of token and  
ReplicaPlacementStrategy

# CONSISTENCY

## WRITE

Consistency Level ONE




**token** - primary replica node  
**replicaPlacementAlgorithm** -  
remaining replica nodes

# CONSISTENCY

## WRITE

Consistency Level ONE

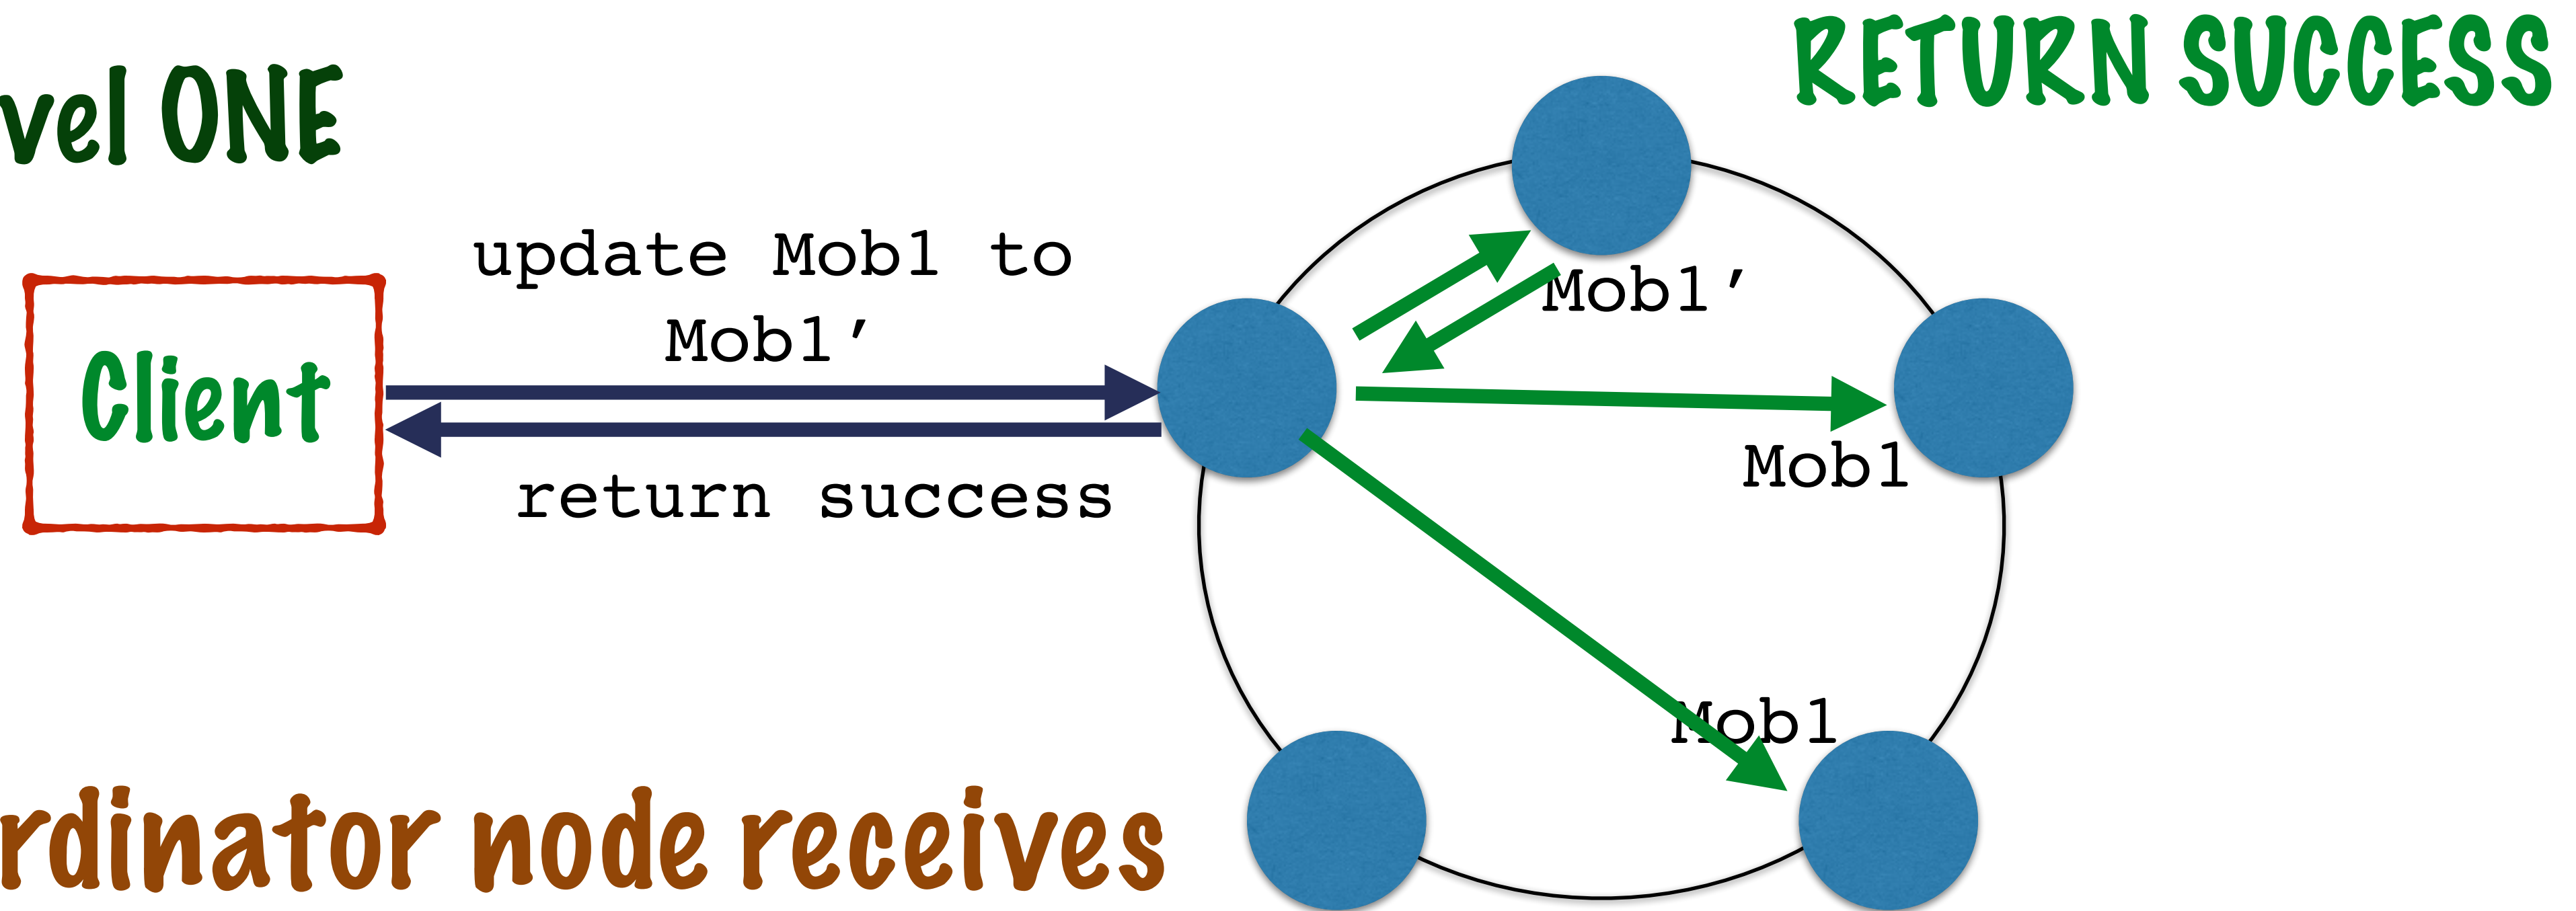


It sends the request  
to all the nodes  
simultaneously

# CONSISTENCY

## WRITE

Consistency Level ONE



As soon as coordinator node receives success response from **ONE NODE**, it returns success to the client



# CONSISTENCY

WRITE

ONE

ALL

QUORUM

LOCAL\_QUORUM

# CONSISTENCY

## WRITE

## ALL

**All** replicas need to be updated and then the write operation returns a success



# CONSISTENCY

WRITE

ONE

ALL

QUORUM

LOCAL\_QUORUM

# CONSISTENCY

## WRITE

# QUORUM

A **minimum** number of replicas (a **quorum**)  
needs to be updated for the write operation  
to return a success

# CONSISTENCY WRITE

Consistency Level QUORUM

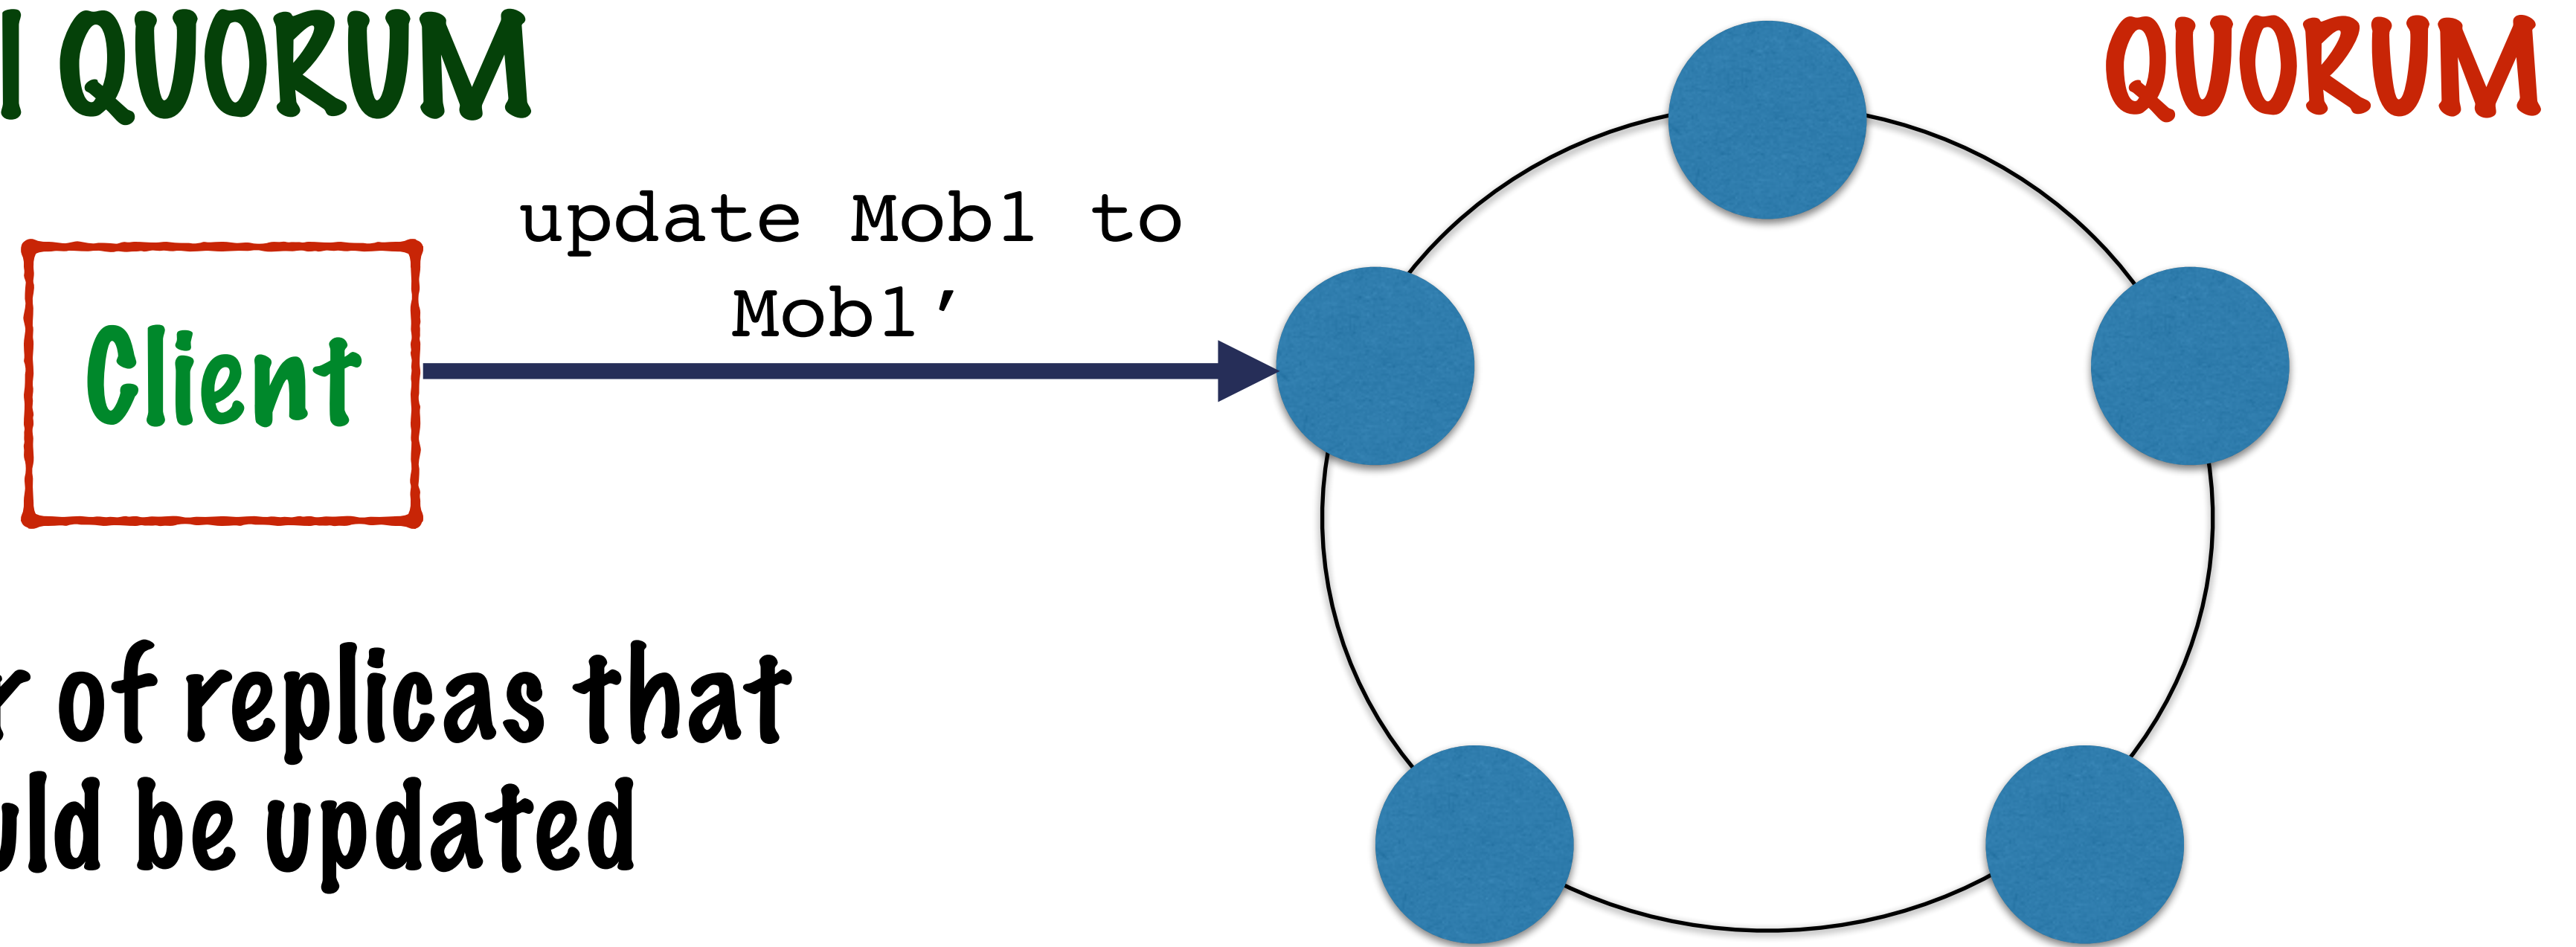
minimum number of  
people  
required to attend the  
meeting so the meeting  
can be conducted

QUORUM



# CONSISTENCY WRITE

Consistency Level QUORUM



Number of replicas that  
should be updated

after which we return  
success to client

# CONSISTENCY


## WRITE

### Consistency Level QUORUM

quorum = 2

Client

update Mob1 to  
Mob1'



Coordinator node determines  
the replica nodes for Mob1

sends the request to all  
the nodes simultaneously



# CONSISTENCY

## WRITE


### Consistency Level QUORUM

quorum = 2

Client

update Mob1 to  
Mob1'

return success



RETURN  
SUCCESS

As soon as coordinator node receives response from the **quorum (2)** nodes, it returns a response to the client

# CONSISTENCY WRITE


Consistency Level QUORUM

quorum = 2

Client

update Mob1 to  
Mob1'

return success



RETURN  
SUCCESS

Consistency level ONE is  
simply a quorum of 1!

# CONSISTENCY

WRITE

ONE

ALL

QUORUM

LOCAL\_QUORUM



# CONSISTENCY

## WRITE

# LOCAL\_QUORUM



A **minimum** number of replicas (a **quorum**)  
needs to be updated **per datacenter** for the  
write operation to return a success

# CONSISTENCY

## WRITE

Consistency level  
**LOCAL\_QUORUM**

Used for multiple  
datacenters




# CONSISTENCY

## WRITE

quorum = 2

Consistency level  
LOCAL\_QUORUM



**SUCCESS** when 2  
replicas are  
updated in each  
datacenter

# CONSISTENCY

## WRITE

ONE

ALL

QUORUM

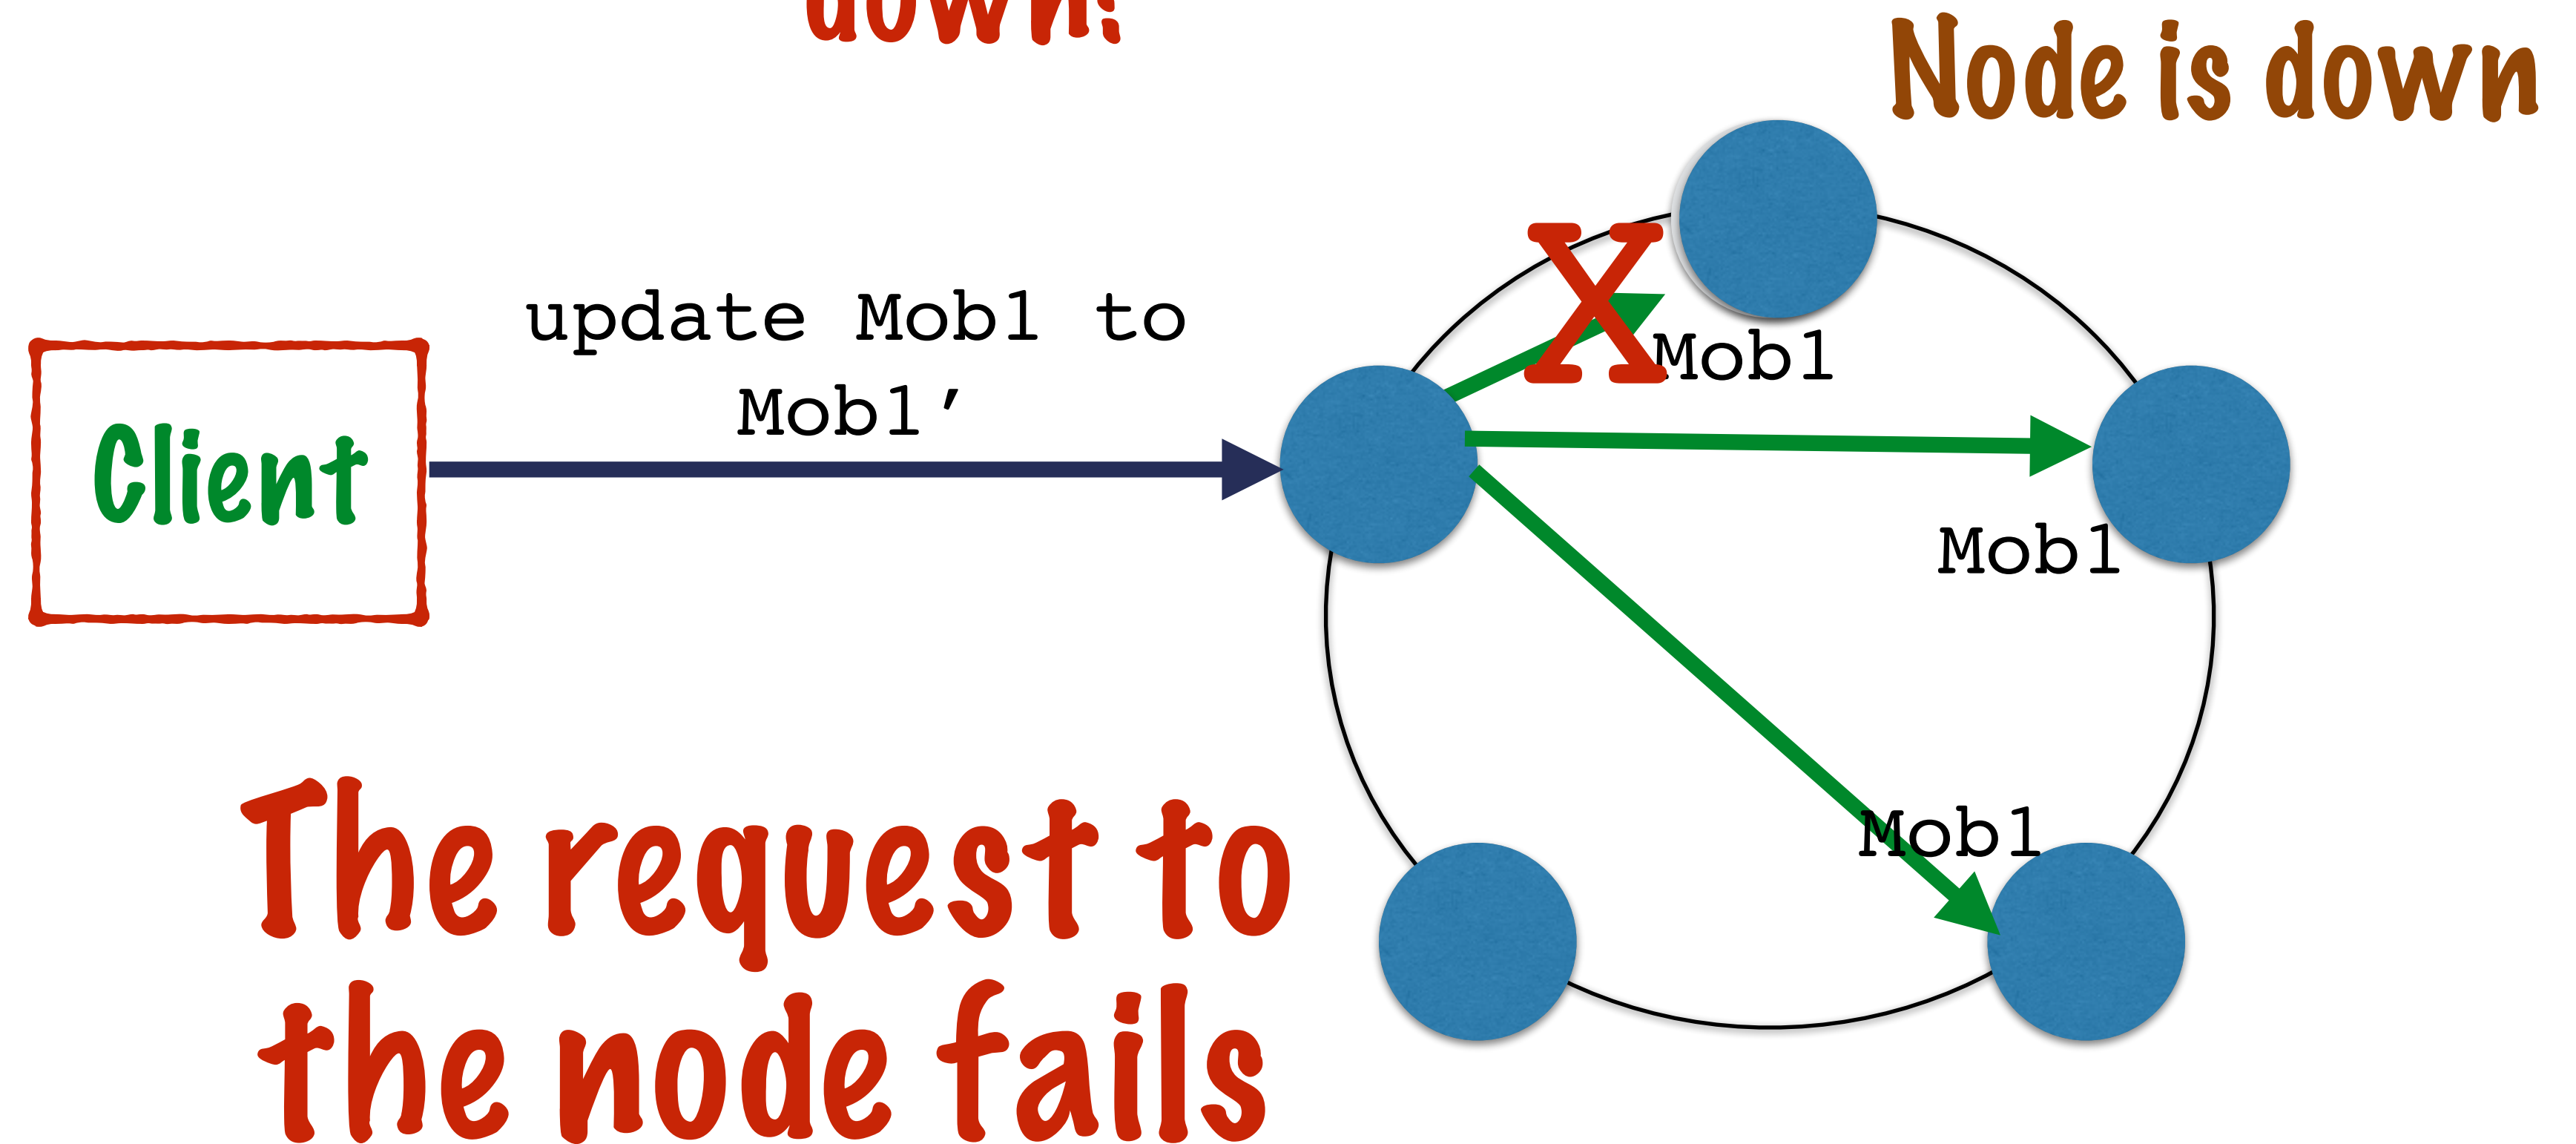
LOCAL\_QUORUM

# CONSISTENCY

**What happens if the  
replica node is down?**

# CONSISTENCY

What happens if the replica node is down?






# CONSISTENCY

What happens if the replica node is down?

Node is down




Coordinator node keeps the data in a file with itself

the file is called hint file

# CONSISTENCY

What happens if the replica node is down?



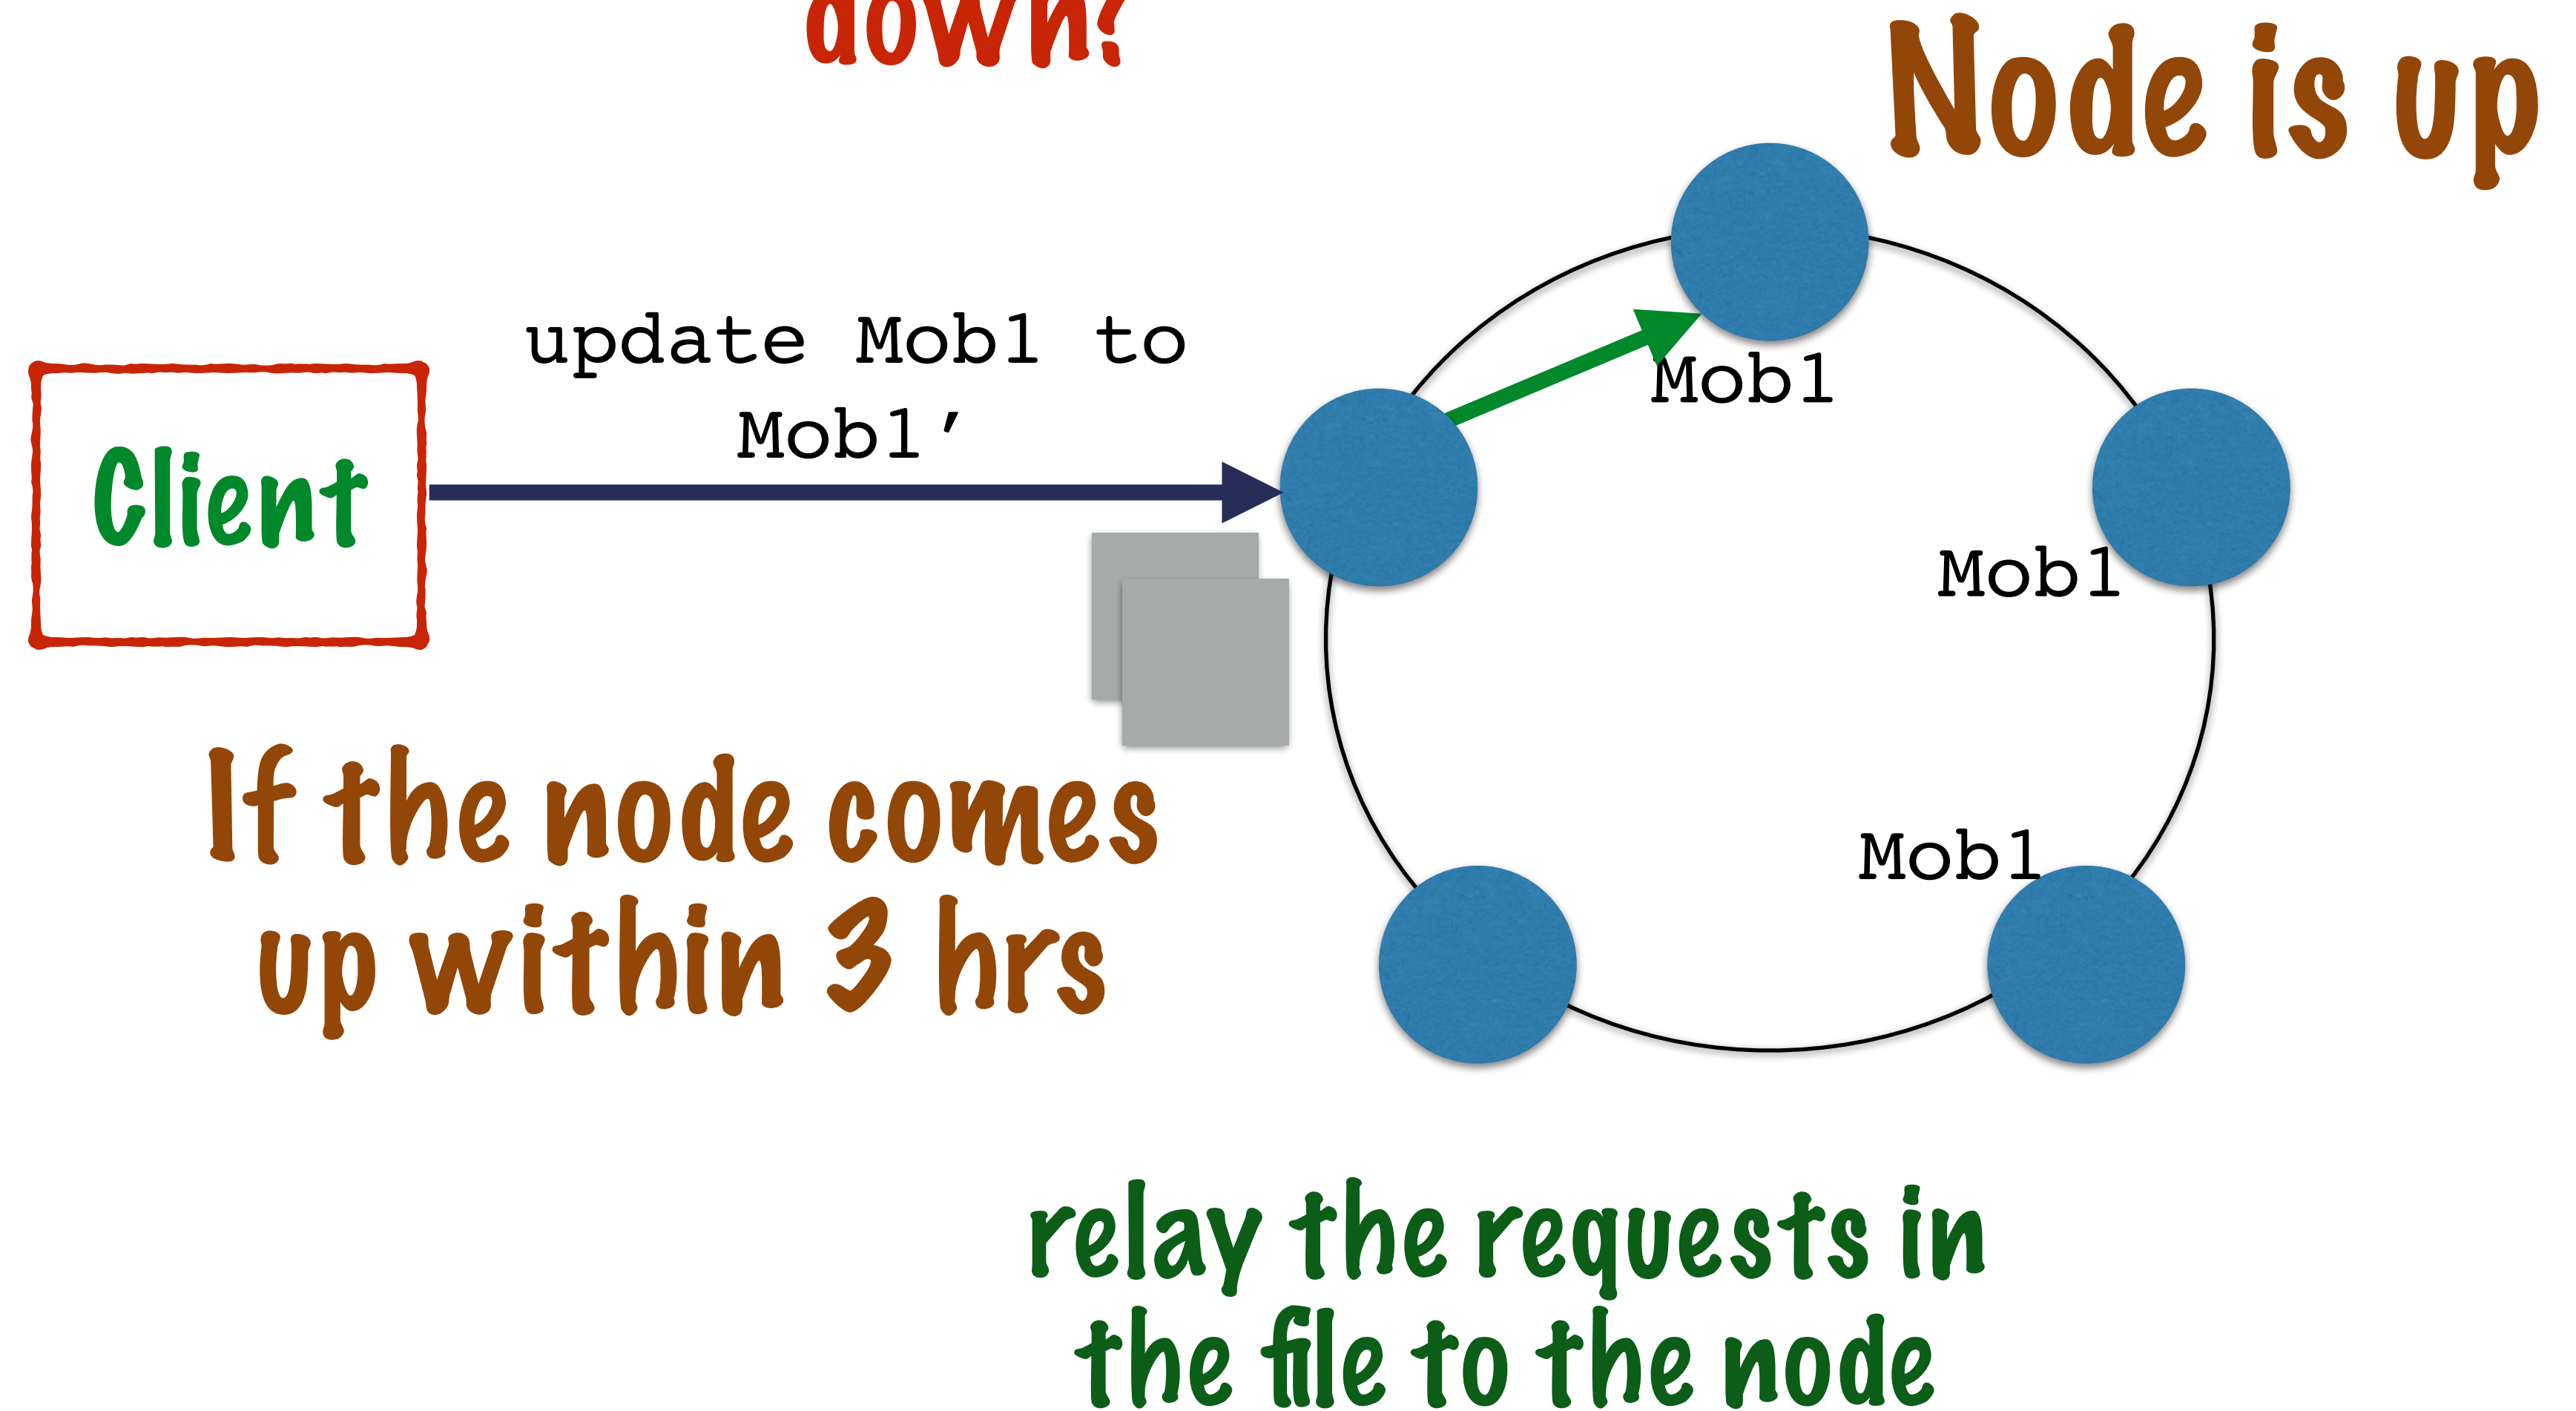
By default it keeps  
this data for 3 hrs

This time interval  
is configurable



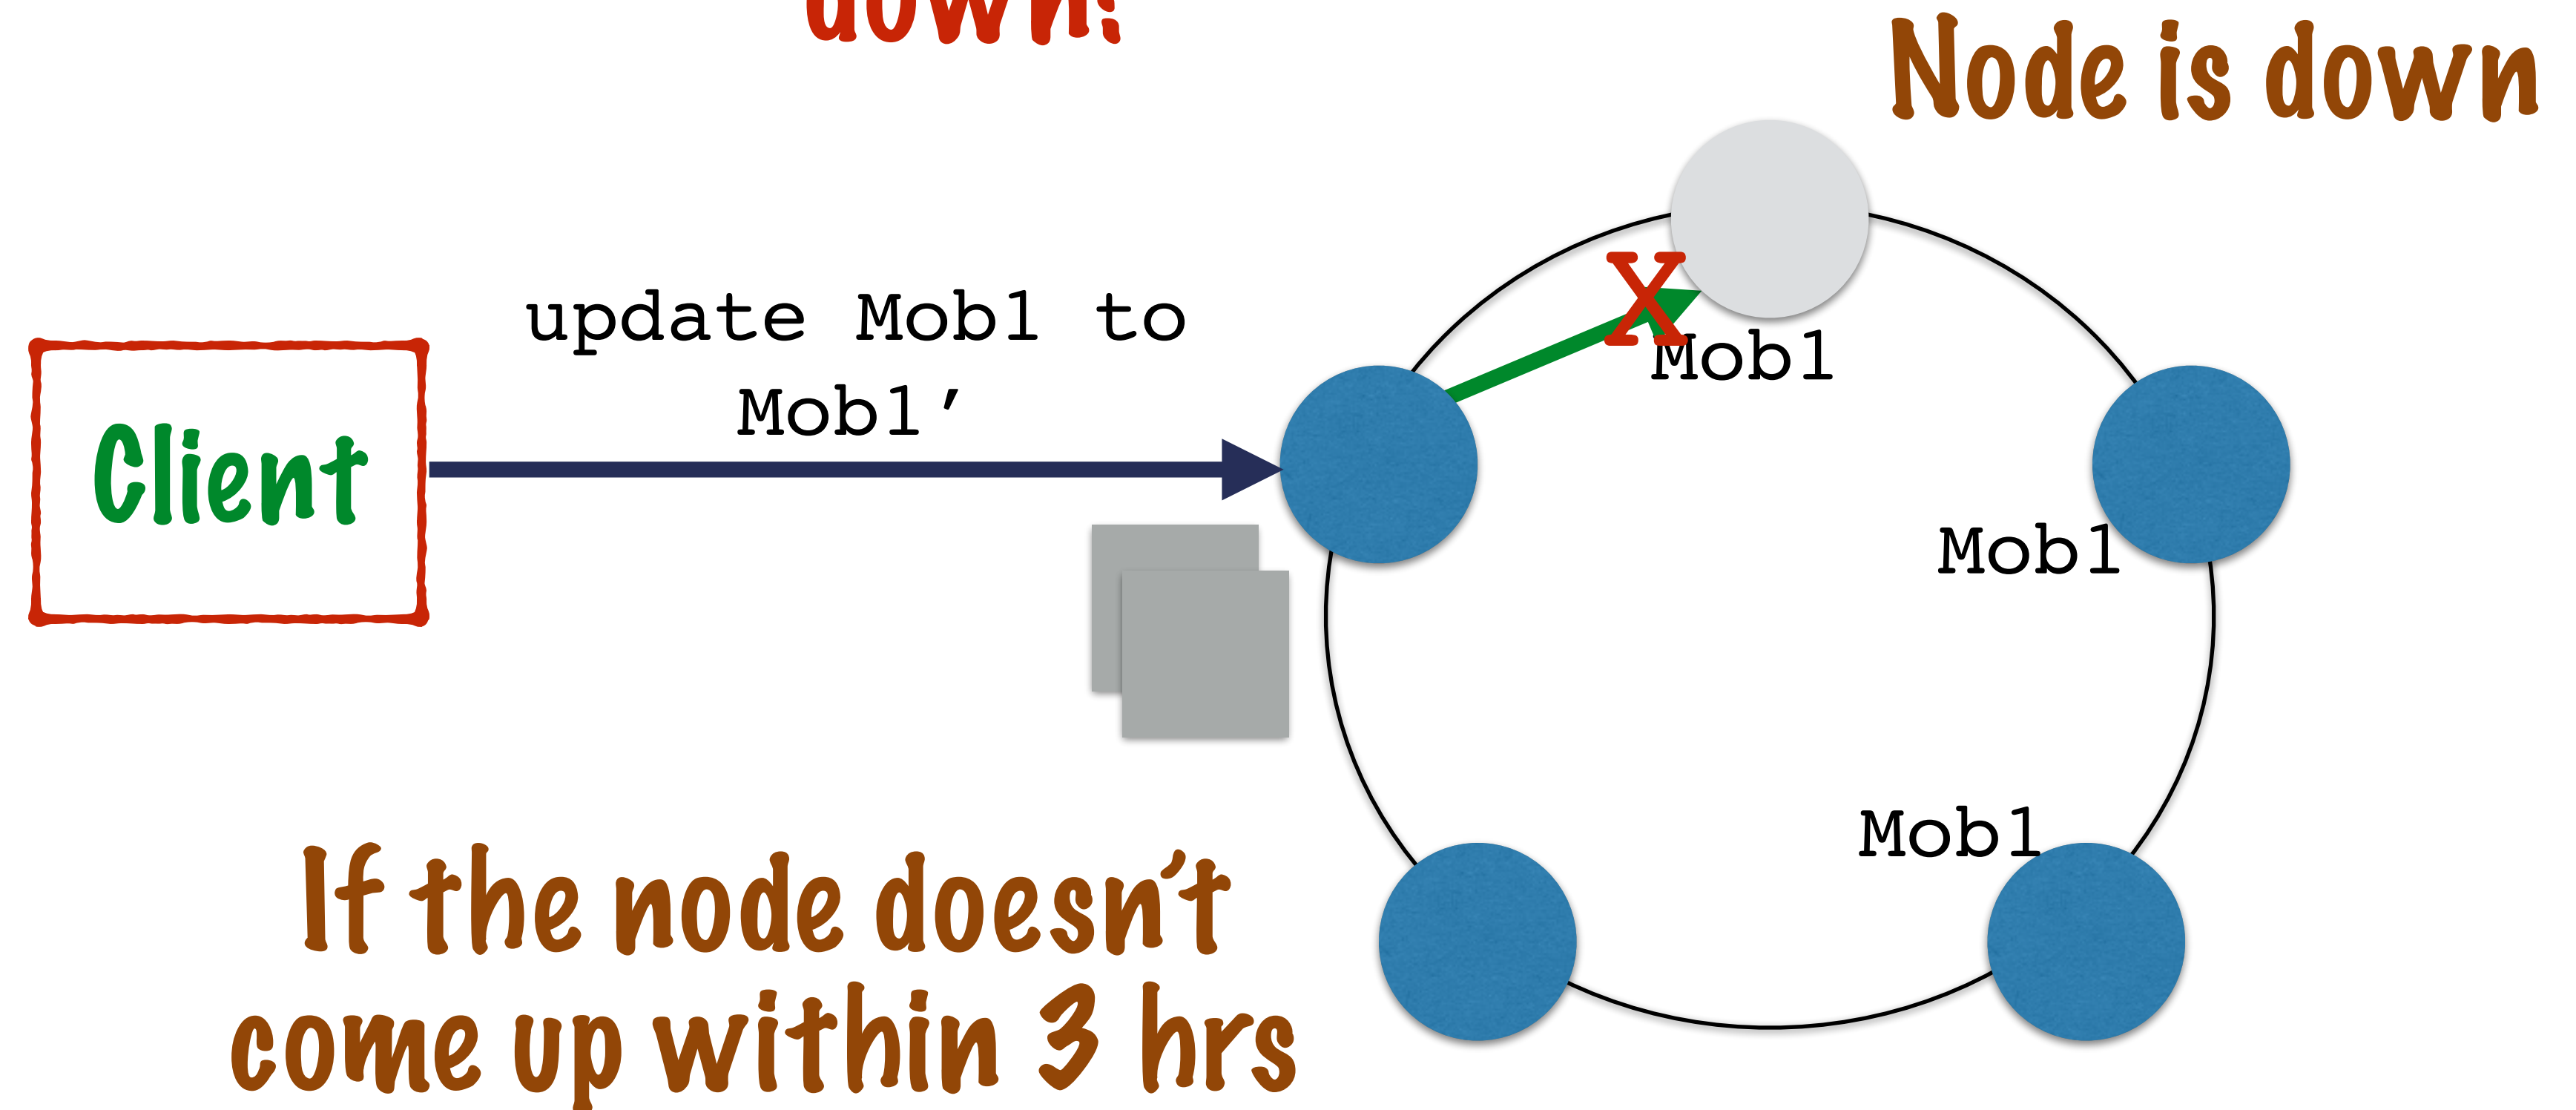
# CONSISTENCY

What happens if the replica node is down?



# CONSISTENCY

What happens if the replica node is down?




If the node doesn't come up within 3 hrs

purge the file

# CONSISTENCY

This mechanism is called  
**HINTED HAND OFF**



It is used to maintain  
consistency

purge the file

Depending on the consistency level specified the request will return SUCCESS or ERROR to the user



Consistency level ALL will return failure for the write request

# CONSISTENCY

Lets see the read  
consistency levels

# CONSISTENCY

READ

ONE

ALL

QUORUM

LOCAL\_QUORUM

# CONSISTENCY

## READ

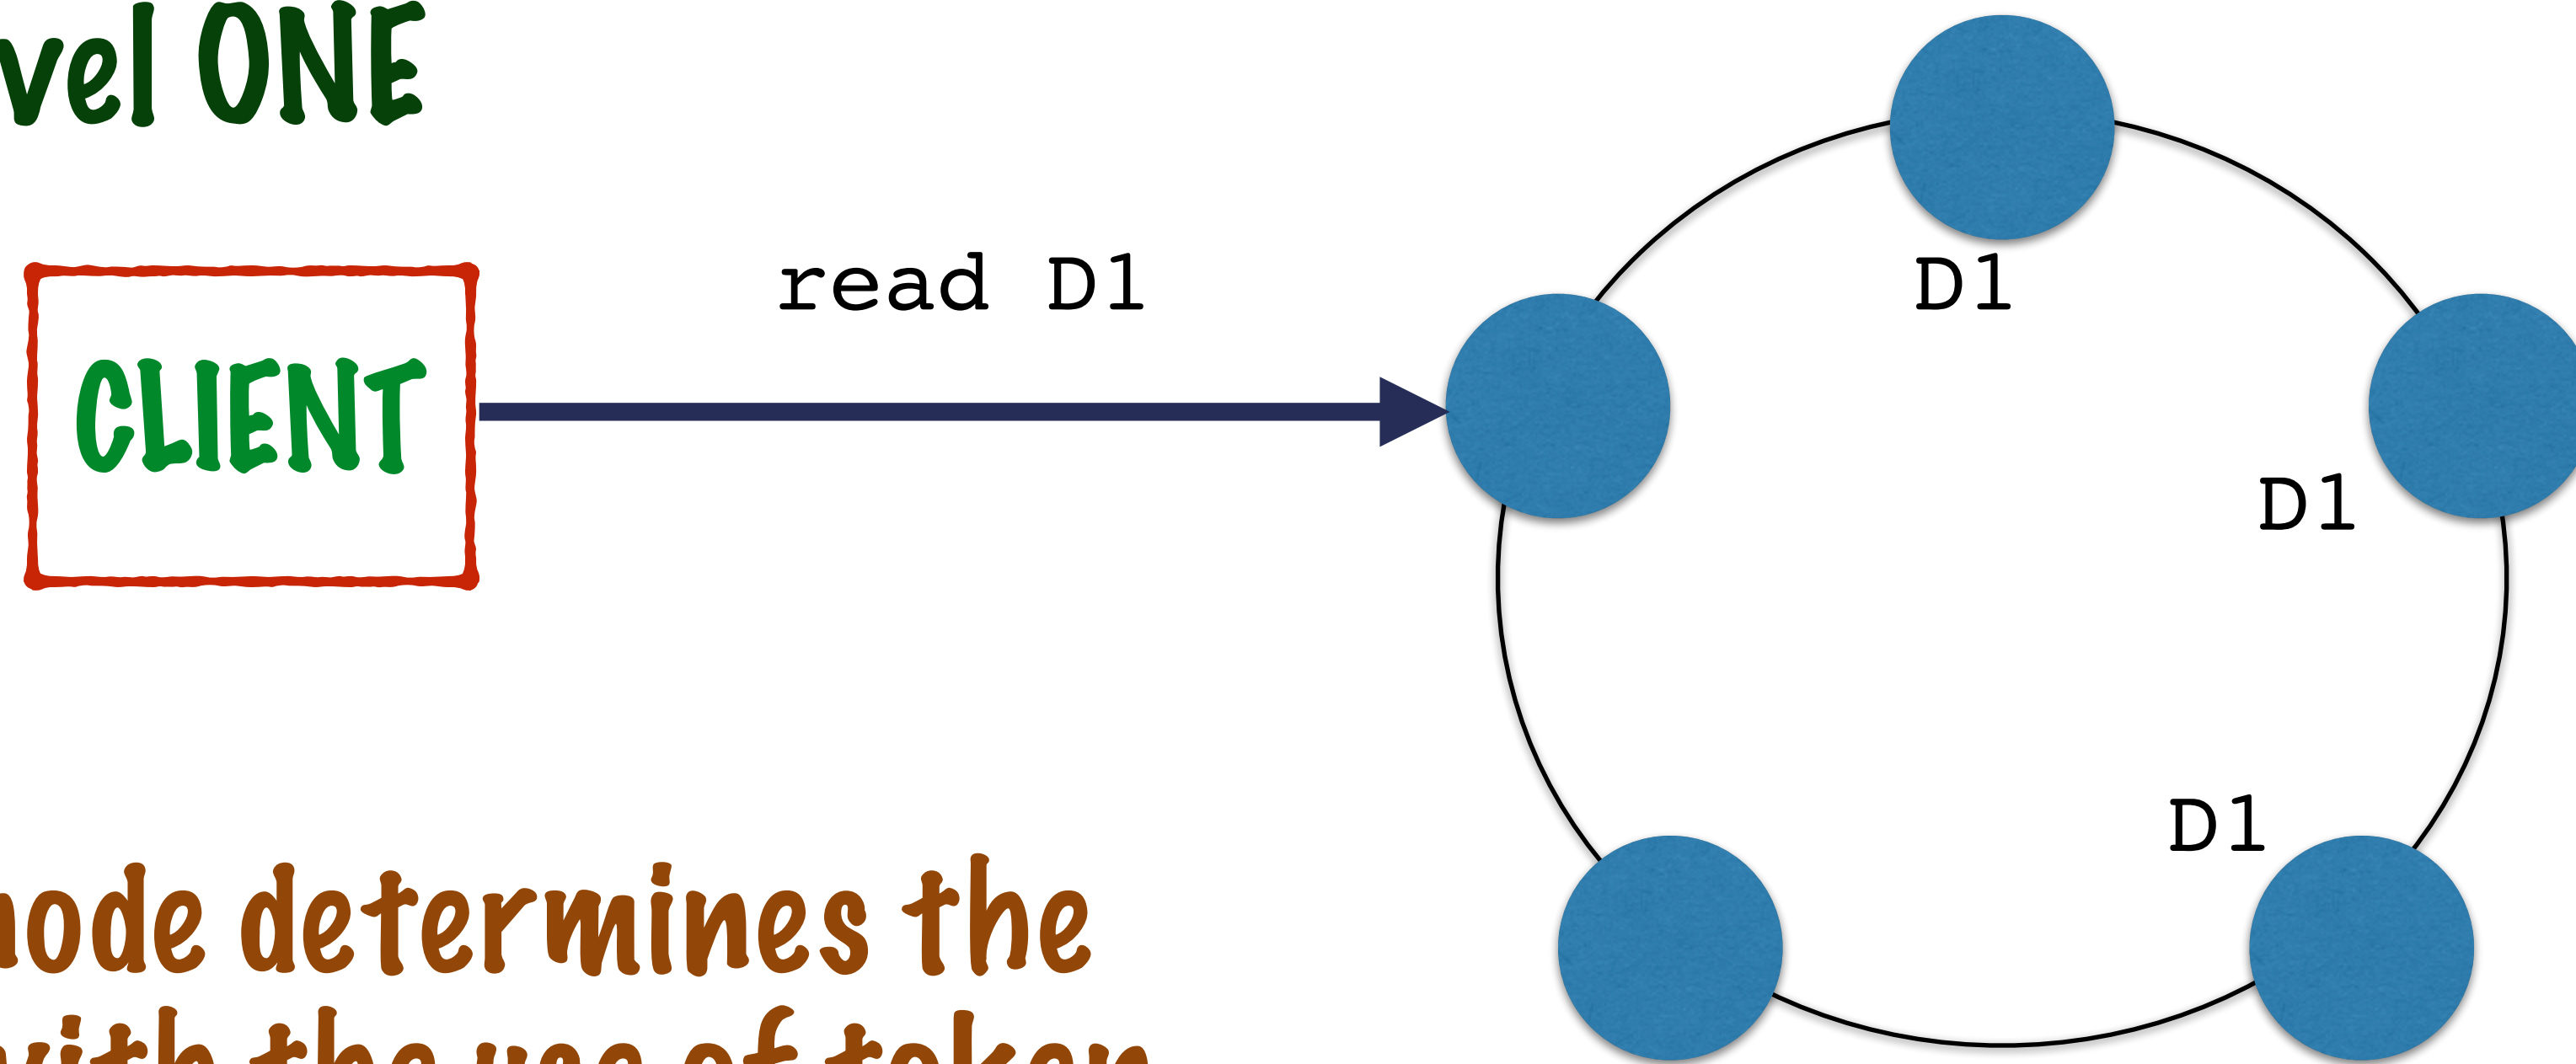
## ONE

**Only one** replica needs to be read and then  
the read operation returns a success



# CONSISTENCY READ

## Consistency Level ONE




coordinator node determines the  
replica nodes with the use of token  
and replicaPlacementAlgorithm



# CONSISTENCY READ

Consistency Level ONE



Now it needs to find the replica  
node which is the fastest

CONSISTENCY

READ

**Snitch** is a program which runs on Cassandra and keeps track of a whole bunch of information about the cluster

**Snitch** determines the  
datacenters and racks  
each node resides in

It also monitors the network latency  
between the nodes and maintains  
this data for each replica

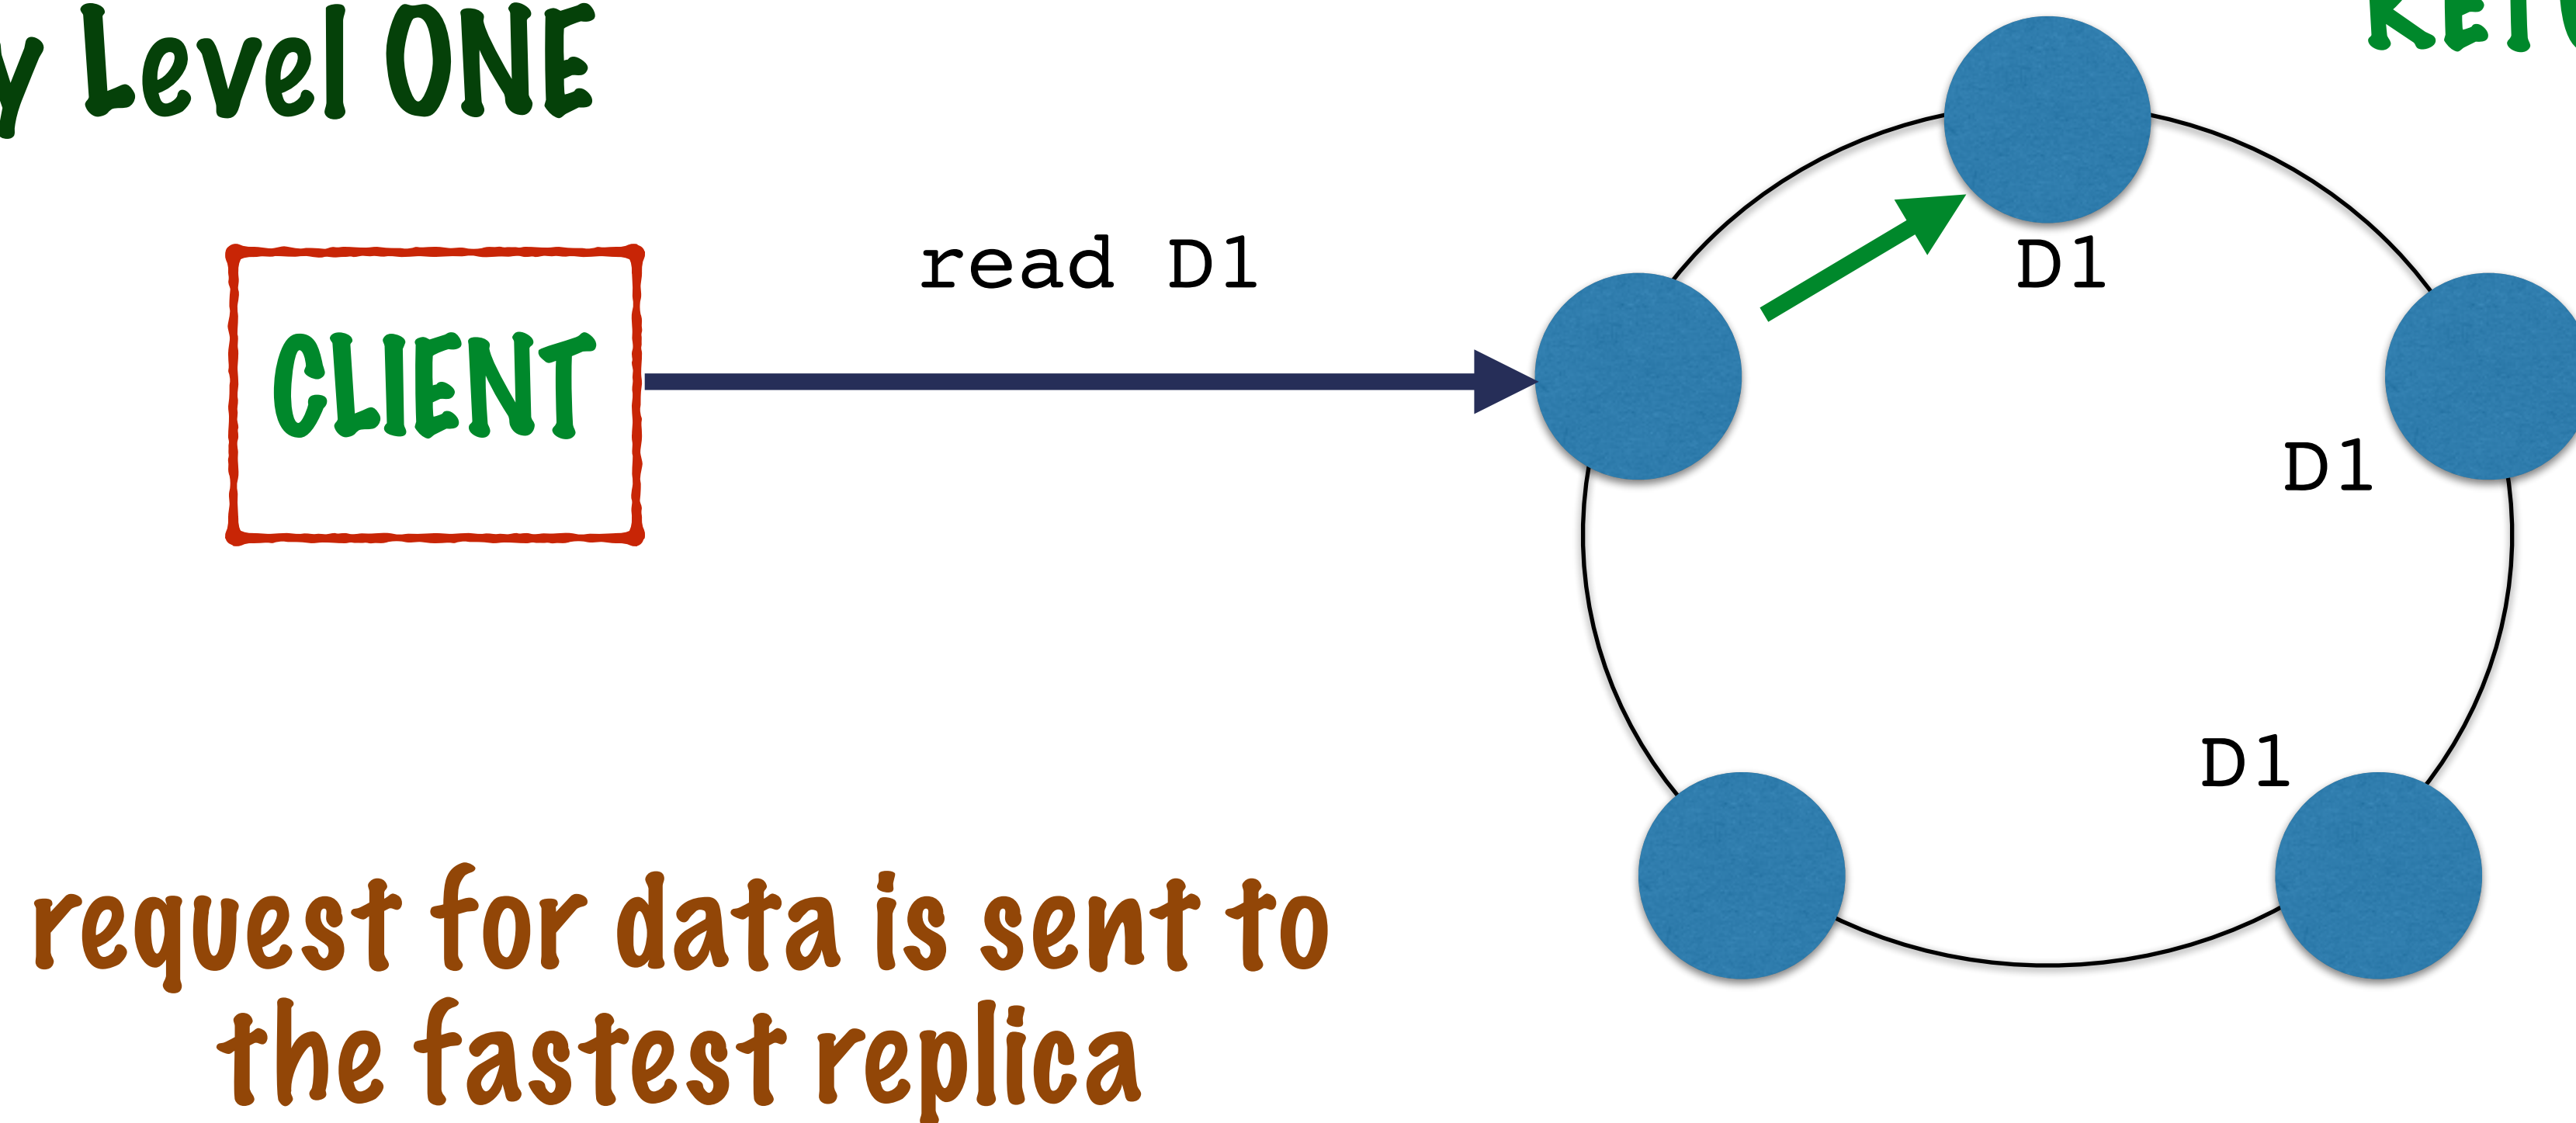
So **snitch** knows which  
is the fastest replica

# CONSISTENCY

## READ

Consistency Level ONE

RETURN DATA

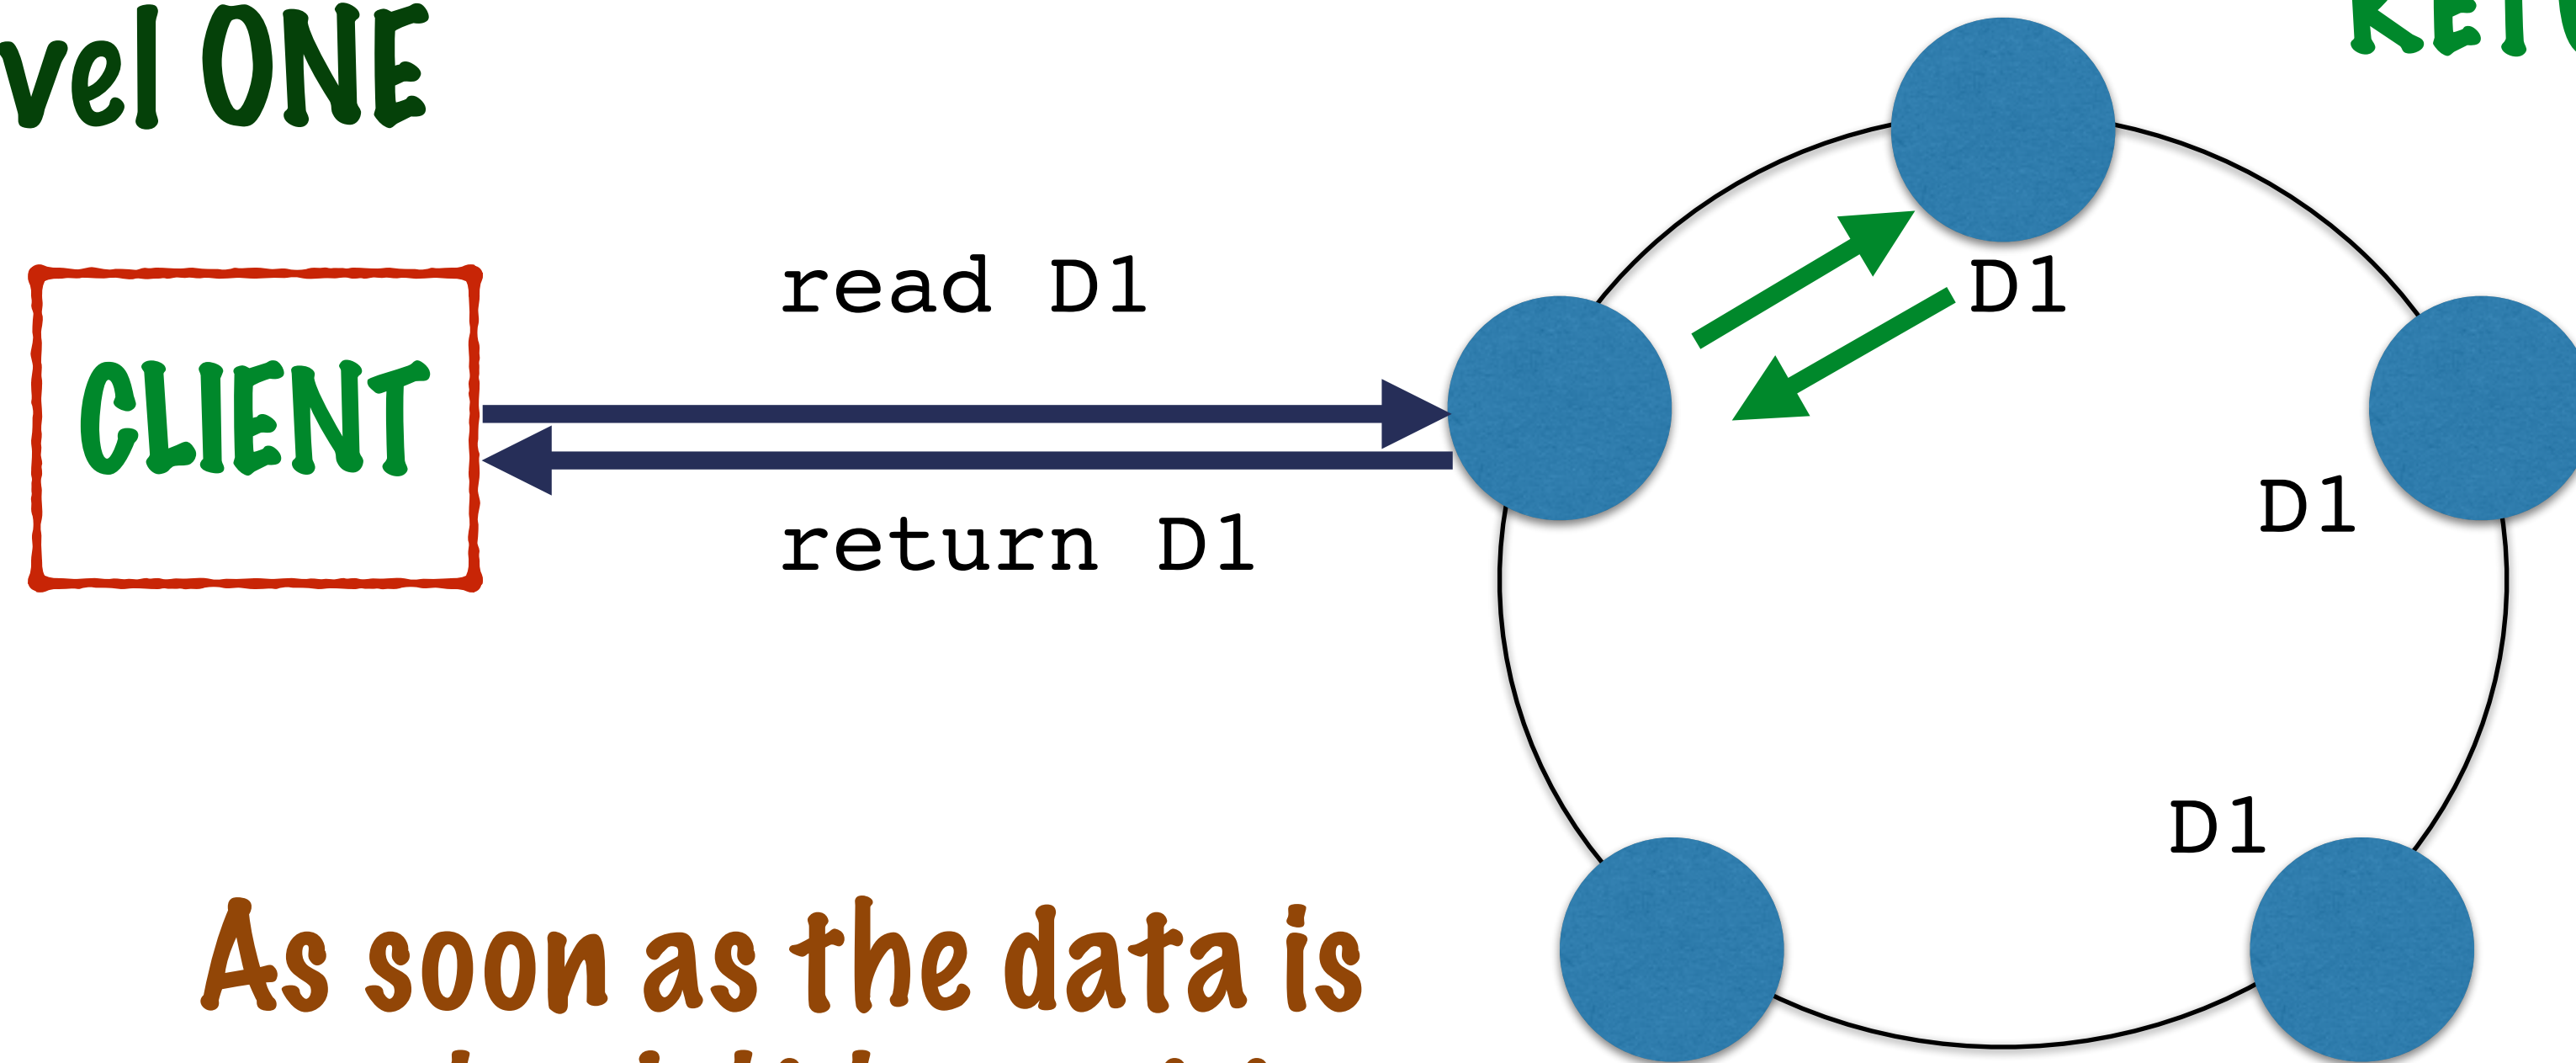


# CONSISTENCY

## READ

Consistency Level ONE

RETURN DATA



As soon as the data is received, it is sent to the client

# CONSISTENCY

READ

ONE

ALL

QUORUM

LOCAL\_QUORUM



# CONSISTENCY

## READ

# QUORUM

A **minimum** number of replicas (a **quorum**)  
needs to be read for the read operation to  
return a success

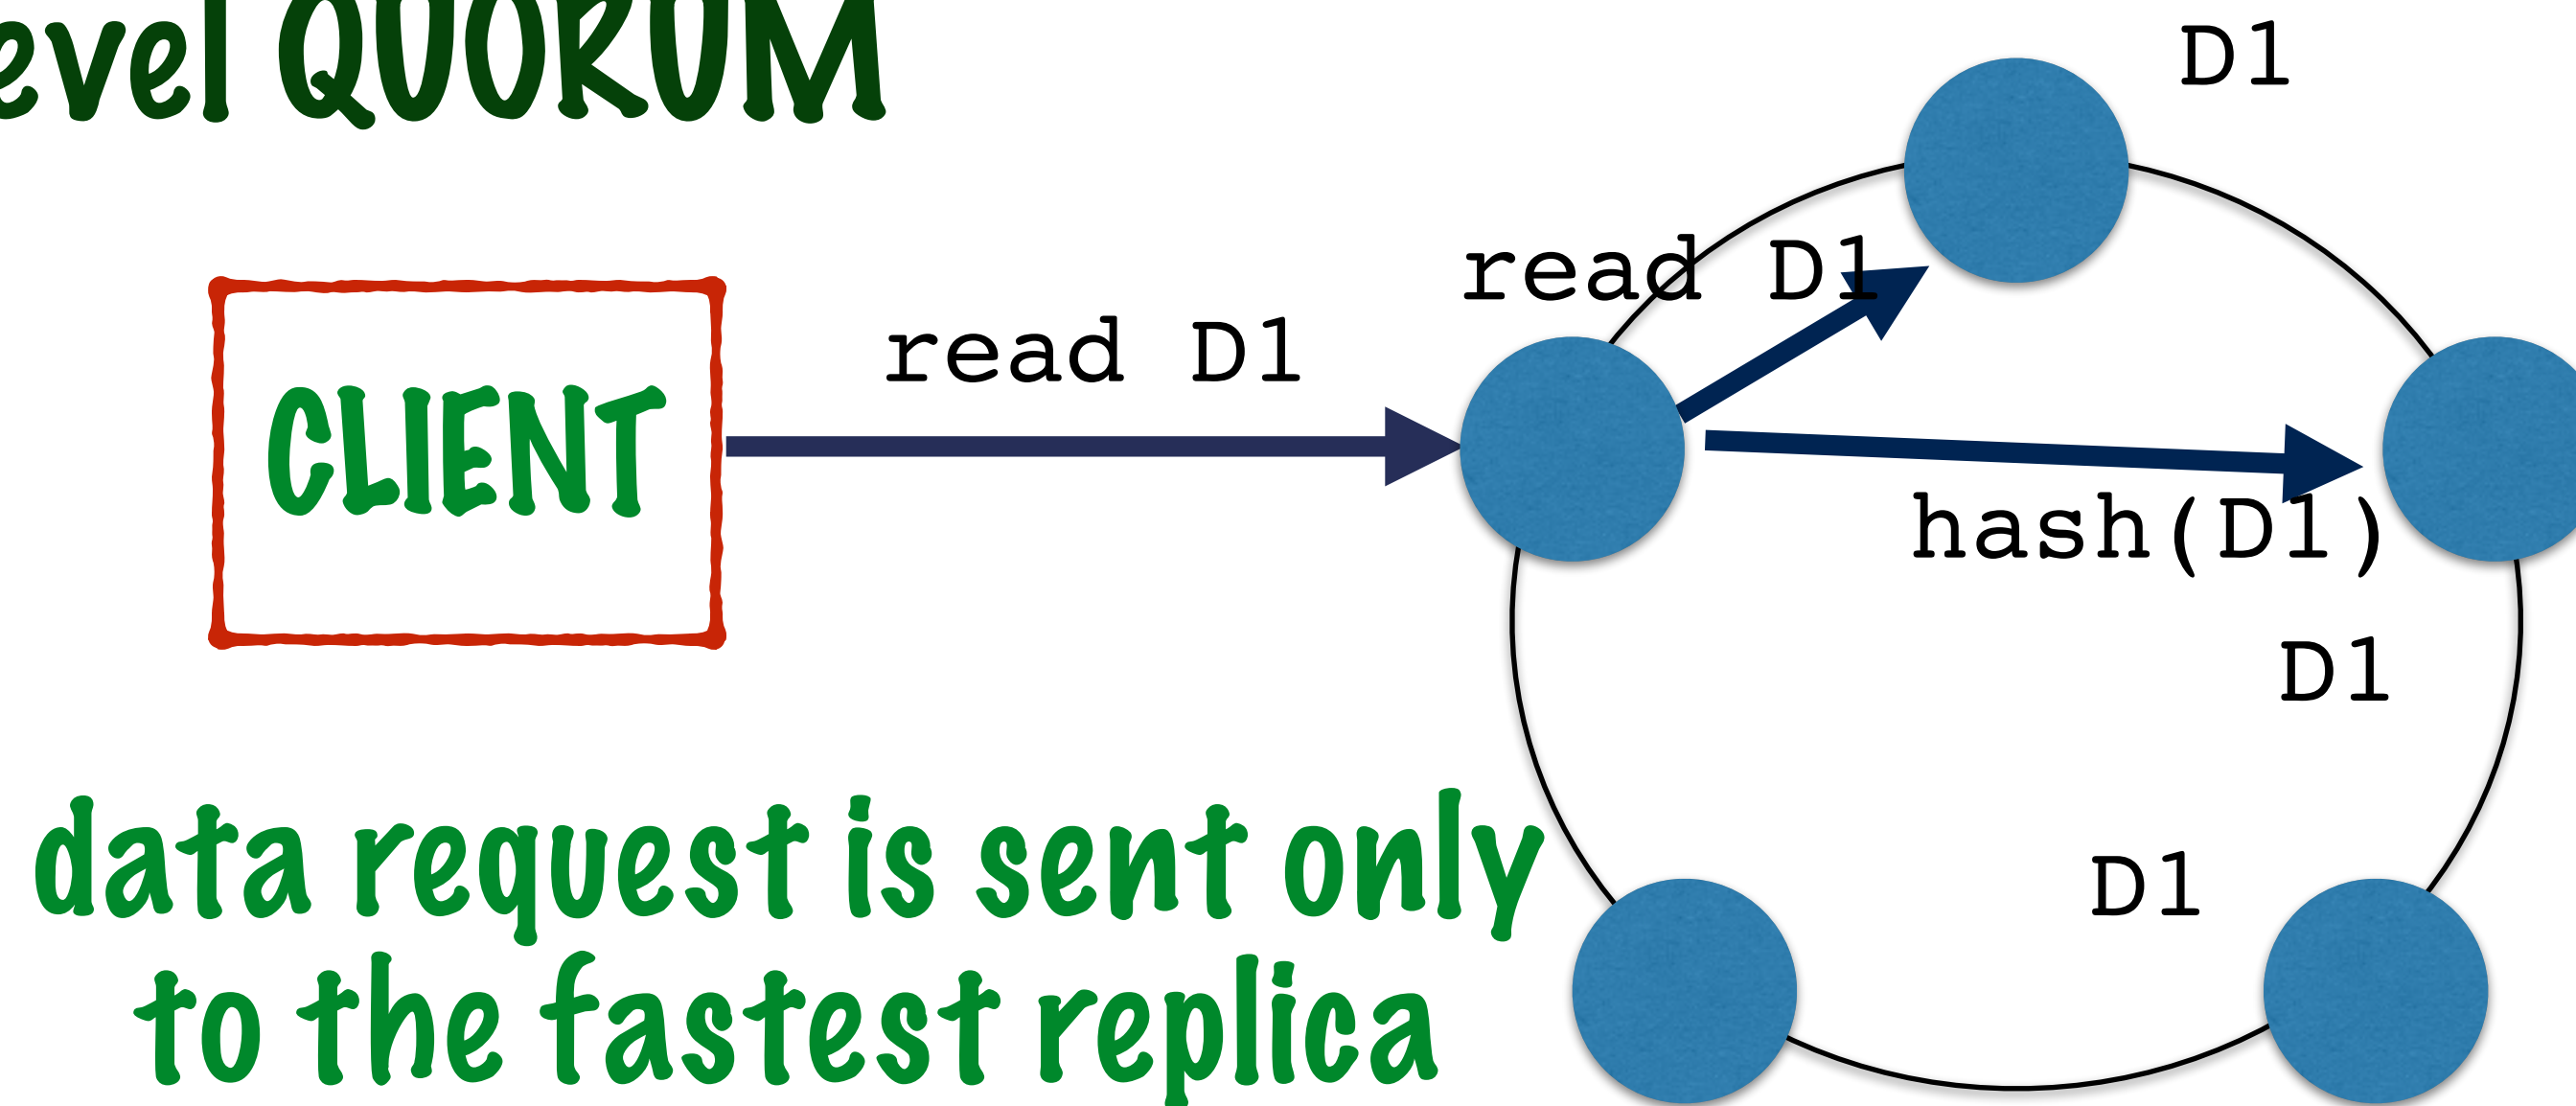


# CONSISTENCY

## READ

quorum = 2

Consistency Level QUORUM



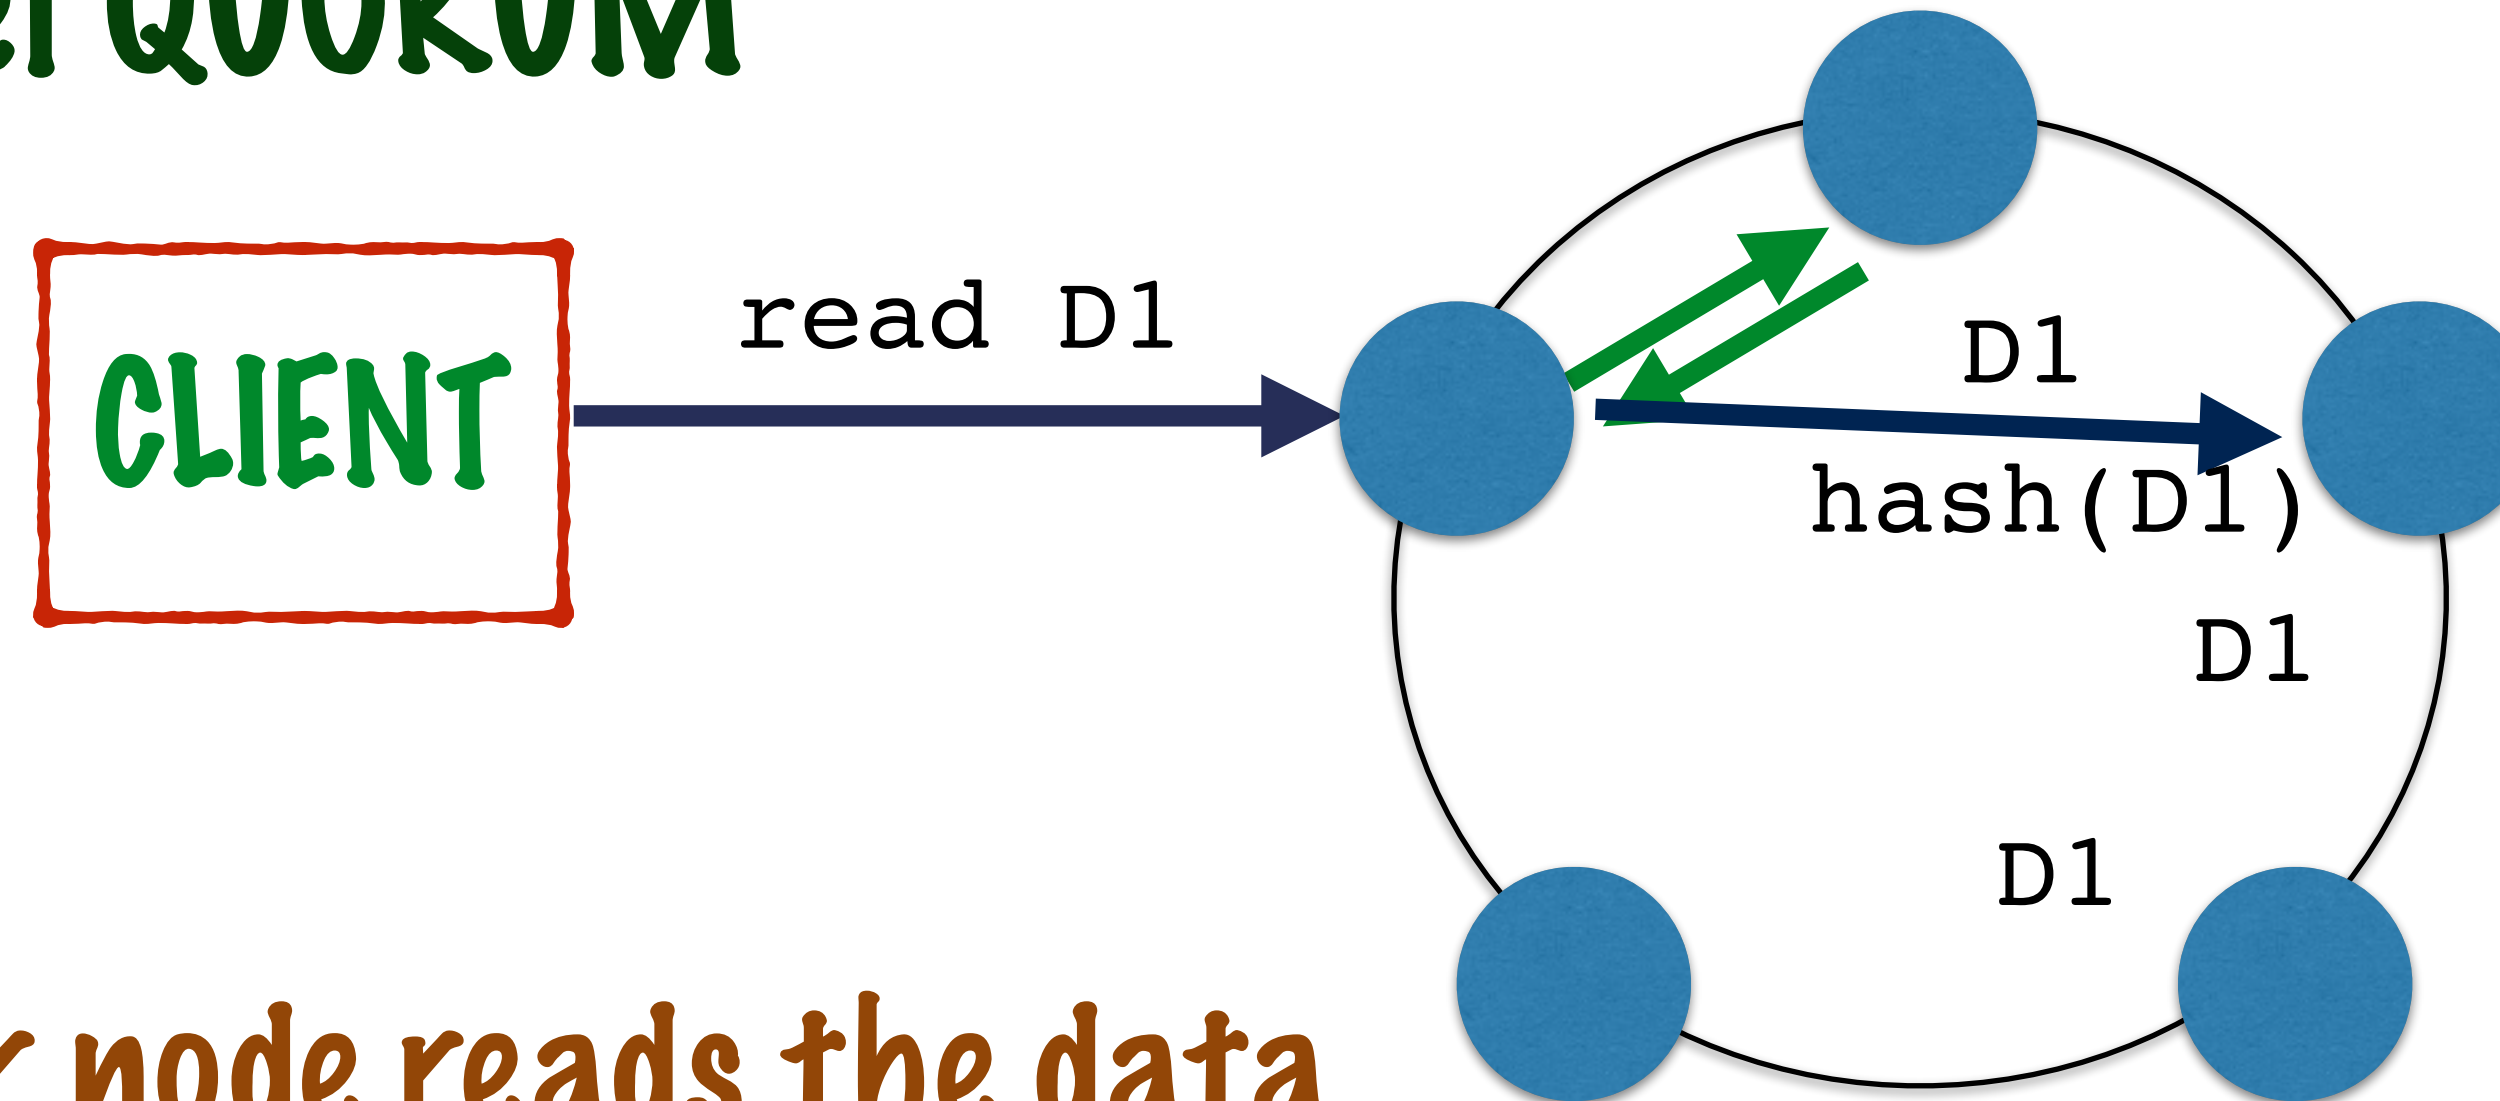
to the other replicas, it sends the request for a **hash** of the data

# CONSISTENCY

## READ

quorum = 2

Consistency Level QUORUM




coordinator node reads the data  
from the first response it receives  
Lets say it is COPY1

# CONSISTENCY

## READ

Consistency Level QUORUM



the other nodes return the  
hash . Lets say it is copy2Hash

# CONSISTENCY


## READ

quorum = 2

Consistency Level QUORUM

CLIENT

read D1



Coordinator node creates  
hash of COPY1 call it  
COPY1HASH

It compares the two hashes


$COPY1HASH == COPY2HASH?$

# CONSISTENCY

## READ

quorum = 2

Consistency Level QUORUM



COPY1HASH == COPY2HASH?

if the hashes **match**, then  
return the data to the client




# CONSISTENCY

READ

quorum = 2

Consistency Level QUORUM



But what if the  
hashes do not match?



# CONSISTENCY

## READ

quorum = 2

Consistency Level QUORUM

CLIENT

read D1

D1


read D1

D1

D1

Coordinator node  
resolves the conflict

It sends a request to 2nd  
node to send the data




# CONSISTENCY

## READ

quorum = 2

### Consistency Level QUORUM



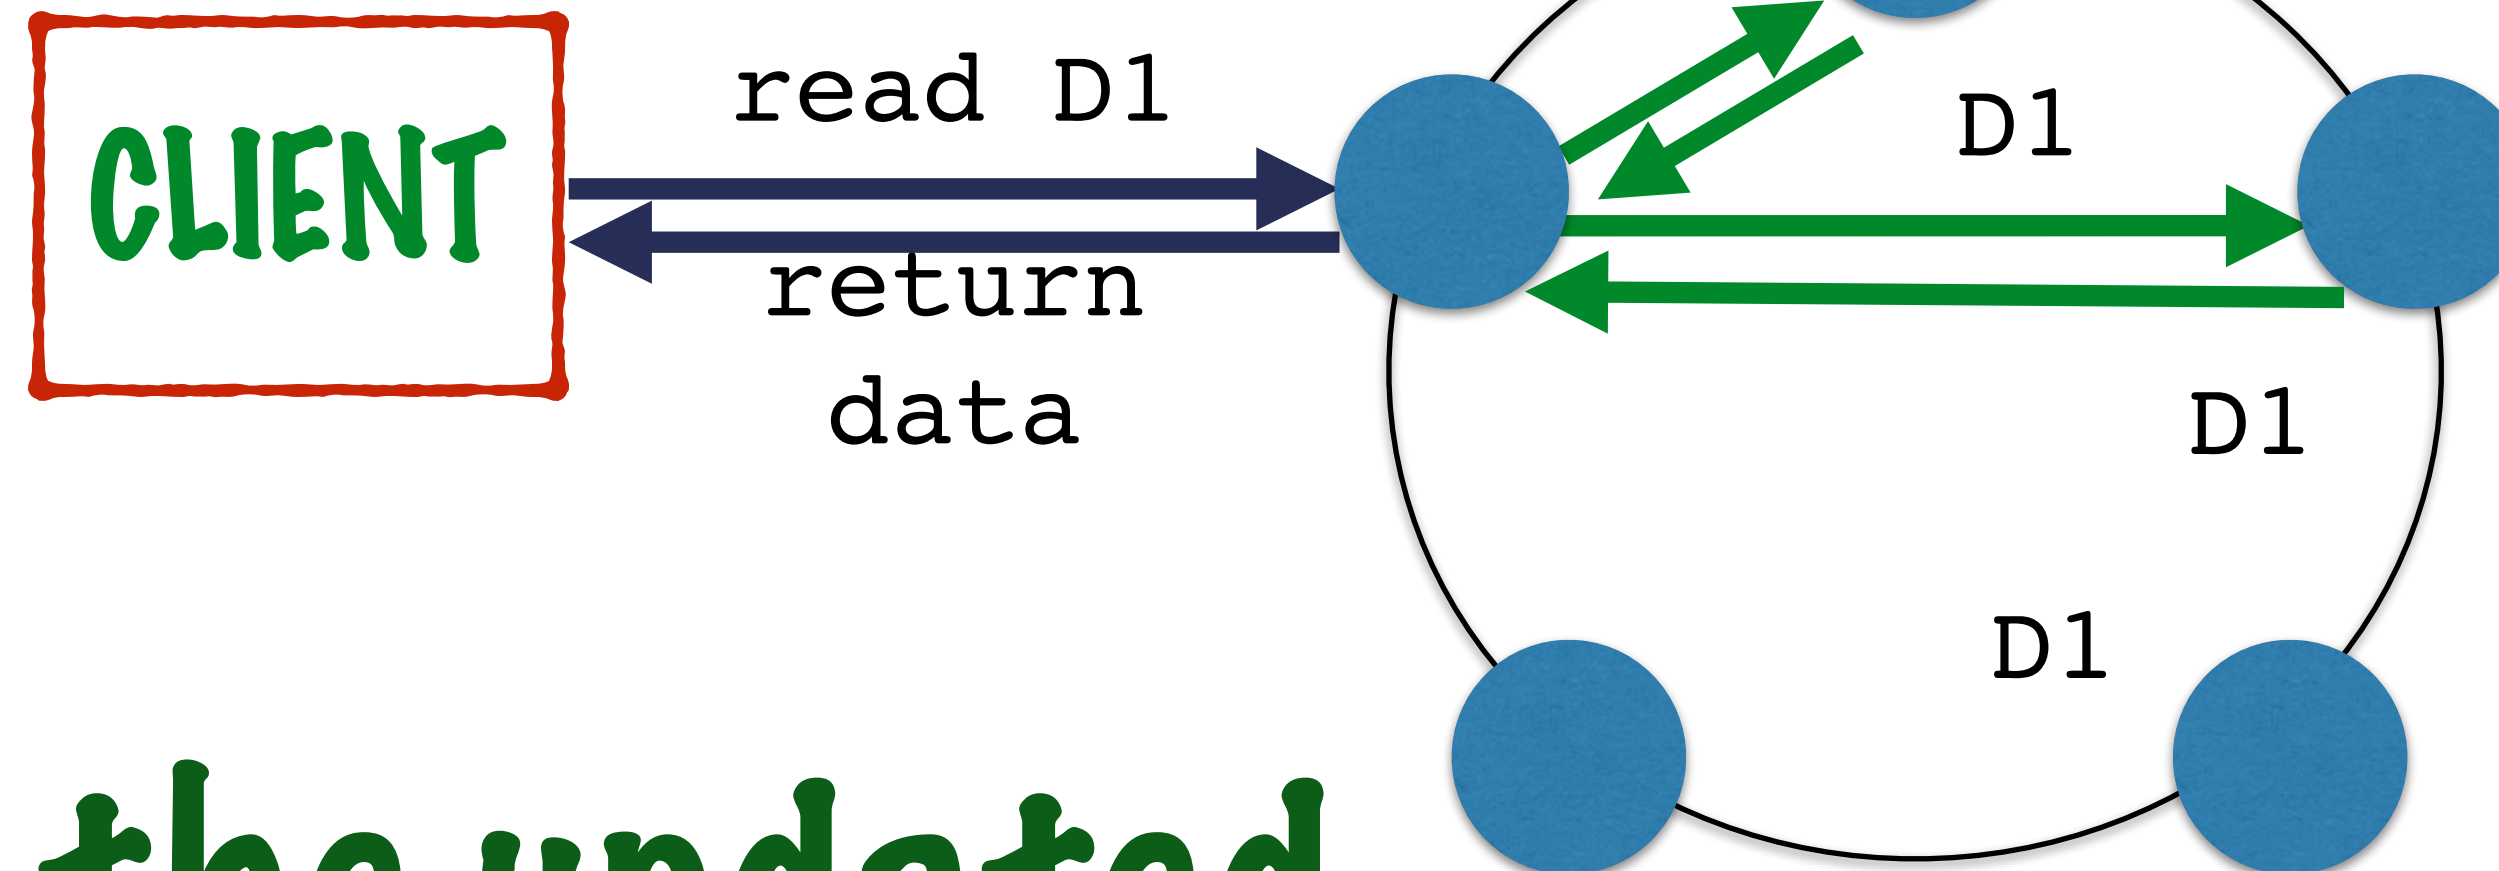
It merges **COPY1** and **COPY2** by keeping the column data with the latest timestamp

# CONSISTENCY

## READ

quorum = 2

Consistency Level QUORUM



It returns the updated data to the client

# CONSISTENCY


## READ

quorum = 2

Consistency Level QUORUM

CLIENT

In the background it sends  
this updated data to all the  
replica nodes asking them  
to update their data




# CONSISTENCY

## READ

quorum = 2

Consistency Level QUORUM

CLIENT



Now in the background it sends this updated data to all the replica nodes asking them to update their data

This mechanism is called **READ REPAIR**




# CONSISTENCY

## READ

quorum = 2

Consistency Level QUORUM

CLIENT



The frequency at which  
Read Repair happens is  
configurable



# CONSISTENCY


## READ

quorum = 2

Consistency Level QUORUM

CLIENT

**read\_repair\_chance** value is  
the probability that read  
repair will be triggered  
when there is a mismatch




# CONSISTENCY READ

Consistency Level QUORUM

Can we use quorum for a multi  
datacenter cluster?

# CONSISTENCY READ

Consistency Level QUORUM



# CONSISTENCY

## READ

Can we use quorum for a multi  
datacenter cluster?

**NO!**

READ OPERATIONS CAN END UP TAKING  
ARBITRARILY LONG IF THEY ARE  
CROSS-DATACENTER READS

# CONSISTENCY

READ

ONE

ALL

QUORUM

LOCAL\_QUORUM

# CONSISTENCY

## WRITE

# LOCAL\_QUORUM


A **minimum** number of replicas (a **quorum**)  
needs to be read **on one datacenter** for the  
read operation to return a success



# CONSISTENCY

## READ

Consistency Level  
LOCAL\_QUORUM



# CONSISTENCY



## READ

quorum = 2

Consistency Level  
LOCAL\_QUORUM

CLIENT

read D1



requests data from 2  
nodes with D1 replicas  
from current datacenter

# CONSISTENCY



## READ

quorum = 2

Consistency Level  
LOCAL\_QUORUM

CLIENT

read D1



AVOIDS INTER-DATACENTER  
COMMUNICATION DELAYS

# CONSISTENCY

READ

ONE

ALL

QUORUM

LOCAL\_QUORUM

# CONSISTENCY

WRITE

ALL

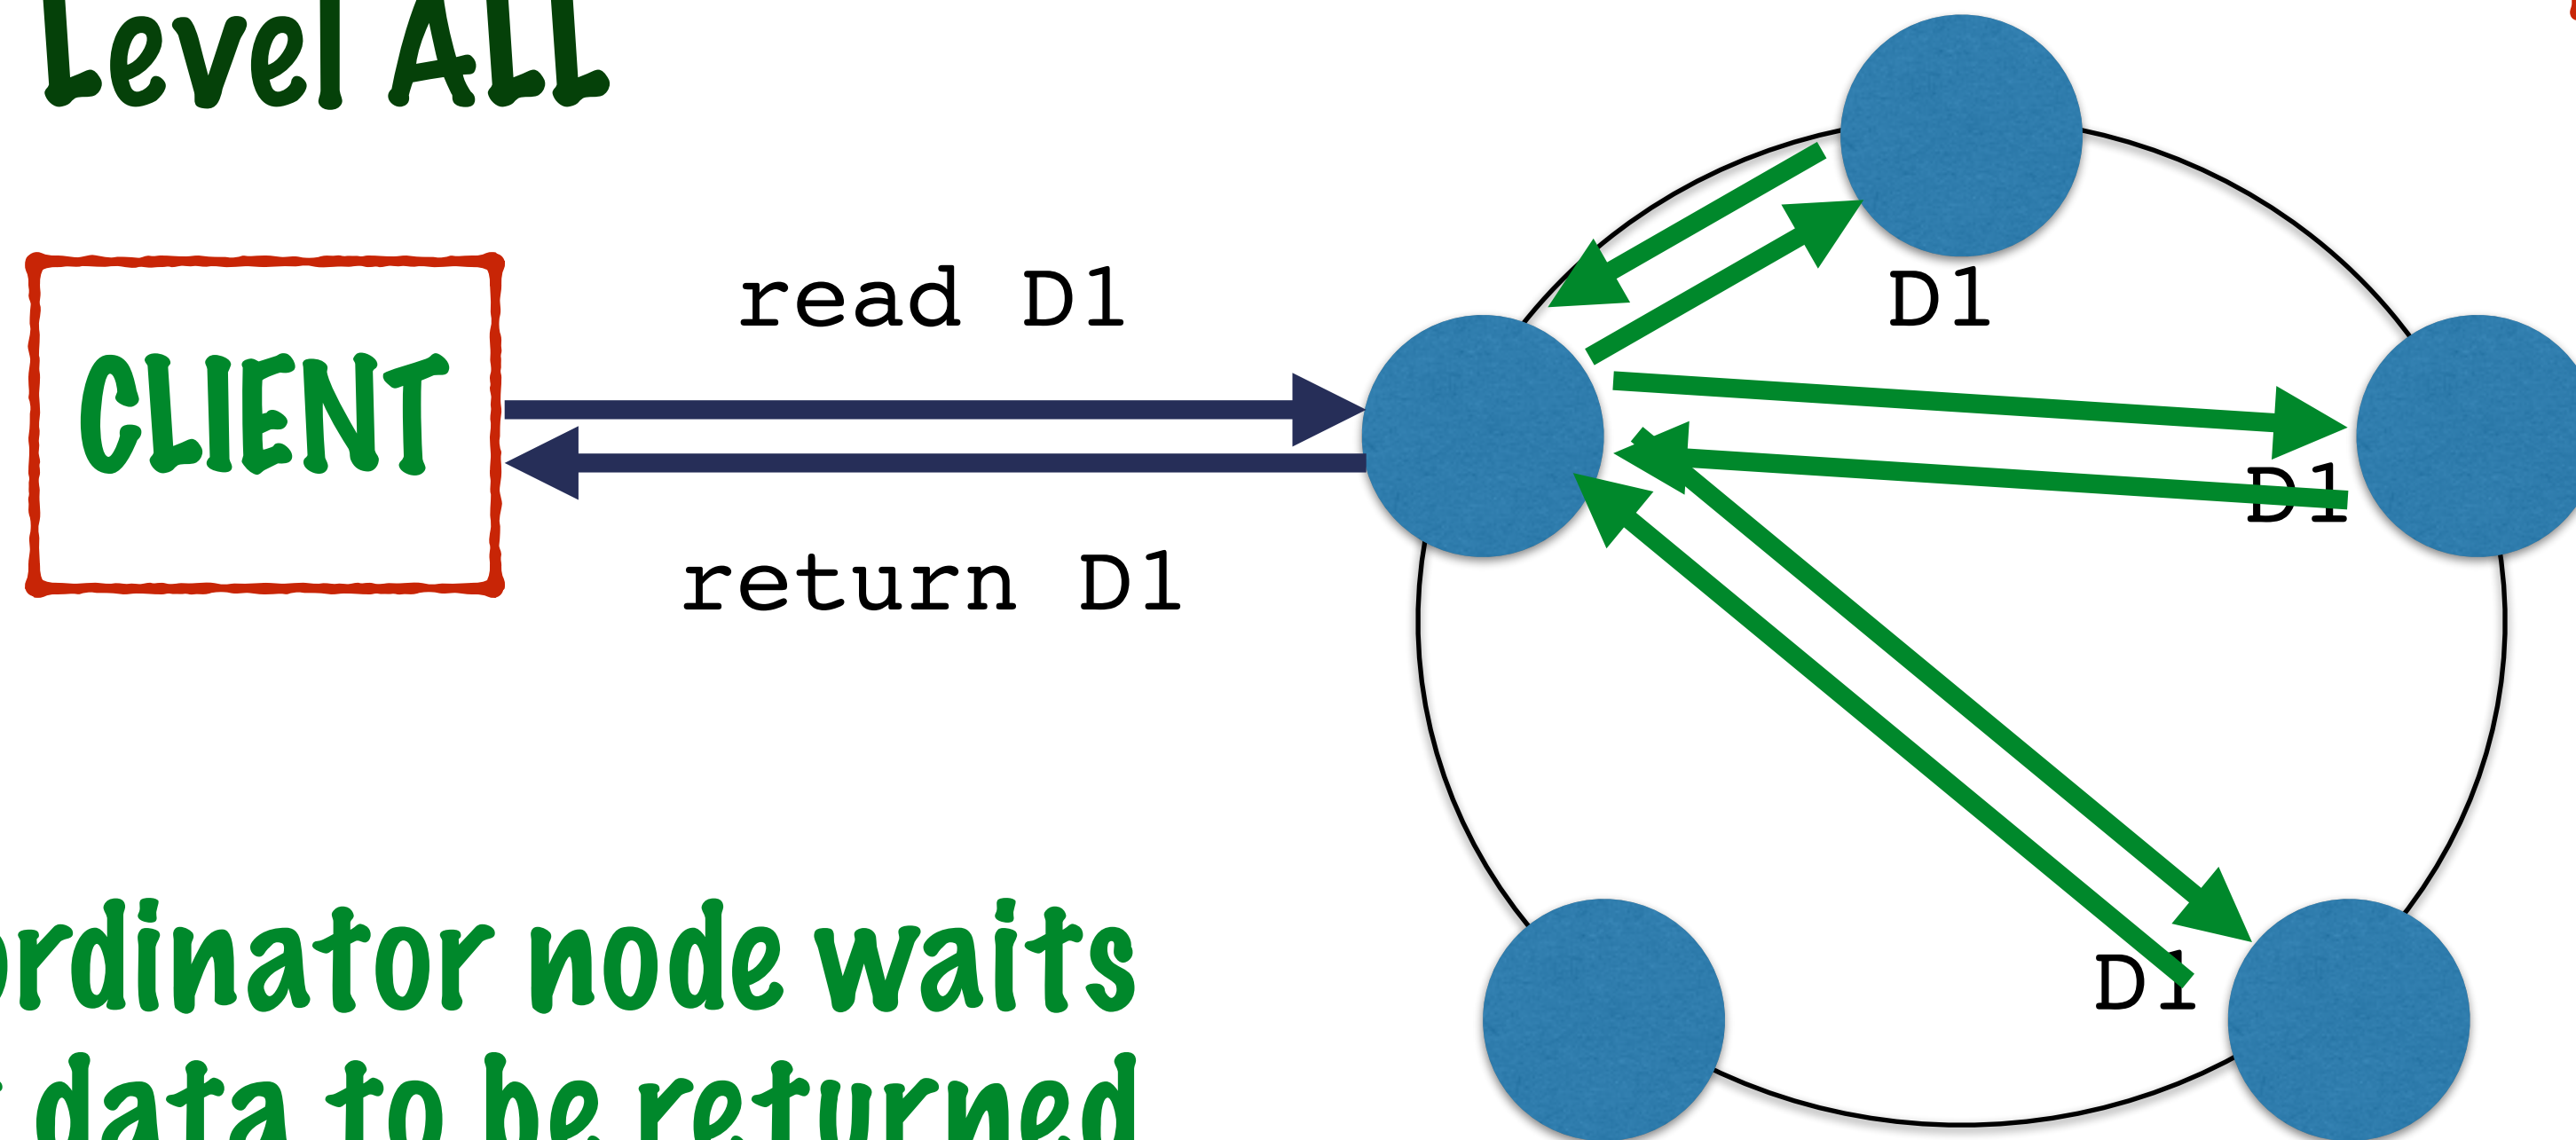
**All** replicas need to be read and only then does the read operation return success

# CONSISTENCY

## READ

request fails if even 1  
node doesn't send the  
response

Consistency Level ALL




coordinator node waits  
for data to be returned  
from all nodes before  
returning to client



# CONSISTENCY

## READ

Consistency Level ALL



STRONGEST  
Consistency

LOWEST  
Availability

# Quorum value is derived from replication\_factor

```
cassandra@cqlsh> CREATE KEYSPACE catalog WITH replication={'class': 'SimpleStrategy',  
                                                         replication_factor': '3'}
```

We can set the replication  
factor of our cluster

# How to determine the replication\_factor (rf)

**rf = 1**

**Only 1 node has data**

**Consistency level = one**

## **PROS**

**faster writes**

**consistent data**

## **CONS**

**disk failure → Loss of data**

**node is down → cluster is unavailable**

# How to determine the replication\_factor (rf)

**rf = 1**


**Only 1 node has data**

**Consistency level = one**

faster writes

consistent data

**PROS**



**CONS**

disk failure  
node is down  
Loss of data  
cluster is  
unavailable

**cons outweigh the pros**

# How to determine the replication\_factor (rf)

$rf = n$       Lets go to the other extreme

consistency level = quorum

## PROS

No data loss at all!

## CONS

Wastage of disk space  
Slower reads and  
writes

NOT ACCEPTABLE

# Relationship between replication factor and quorum

`rf = replication_factor (configurable)`

This is per-datacenter

`quorum = ceiling((sum of all rf + 1)/2)`

Sum across all datacenters

`(read_quorum + write_quorum) > rf`

ensures consistency



# Ensuring consistency

`(read_quorum + write_quorum) > rf`

`write_quorum = rf,`  
`read_quorum = 1`

write succeeds only when  
data is updated on all replicas

read can be from any one  
replica

# Ensuring consistency

`(read_quorum + write_quorum) > rf`

`write_quorum = 1,  
read_quorum = rf`

write succeeds when data is updated  
on 1 replica

During read, all the replica nodes  
are checked for data

# Ensuring consistency

$$(\text{read\_quorum} + \text{write\_quorum}) > rf$$

$$\text{write\_quorum} = \text{read\_quorum} = (rf/2)+1$$

during read operation, there  
would be atleast 1 replica node  
where write has been  
successful

# How to determine the replication\_factor (rf)

For a 10 node cluster

3 replicas of data ensures fault tolerance

$rf = 3$

read consistency = quorum (value : 2),

write consistency = quorum (value : 2)

Cluster will be still available, if  
1 node has failed

$read\_quorum + write\_quorum = 4 > rf$